

Lab Week 1: Example responses

Code ▾

1 R programming

1.1 File I/O

Reading a dataset

- Download the `Cereal.csv` file from the Canvas page and use the `read.csv` command to read in the csv file into R and assign it to the object called `cereal`.

Solution

▼ Code

```
library(readr)
library(tidyverse)
cereal <- read.csv("Cereal.csv", header = TRUE)
###cereal_df # in HTML file, prints out full data frame -
cereal_tbl<- read_csv("Cereal.csv")
cereal_tbl # in HTML file, only prints out first 10 rows -
```

A tibble: 77 × 16

	name	mfr	type	calor... ¹	protein	fat	sodium	fiber
	carbo	sugars	potass					
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
	<dbl>	<dbl>	<dbl>					
1	100%_Bran	N	C	70	4	1	130	10
5	6	280						
2	100%_Natu...	Q	C	120	3	5	15	2
8	8	135						
3	All-Bran	K	C	70	4	1	260	9
7	5	320						
4	All-Bran_...	K	C	50	4	0	140	14
8	0	330						
5	Almond_De...	R	C	110	2	2	200	1
14	8	-1						
6	Apple_Cin...	G	C	110	2	2	180	1.5
10.5	10	70						
7	Apple_Jac...	K	C	110	2	0	125	1
11	14	30						

```

      8 Basic_4      G      C      130      3      2      210      2
18      8      100
      9 Bran_Chex   R      C      90      2      1      200      4
15      6      125
10 Bran_Flak... P      C      90      3      0      210      5
13      5      190
# ... with 67 more rows, 5 more variables: vitamins <dbl>, shelf
<dbl>,
#   weight <dbl>, cups <dbl>, rating <dbl>, and abbreviated
variable name
#   ^calories

```

1.2 Data Types

Data frames

- There should be a default dataset in R called `cereal`. Use the `head` function to inspect the first few lines of the data frame and use `class` to check that `cereal` is in fact a data frame.

Solution

▼ Code

```

#### Base R
head(cereal)

```

```

              name mfr type calories protein fat
sodium fiber carbo
1              100%_Bran   N    C        70        4    1
130  10.0    5.0
2              100%_Natural_Bran   Q    C        120        3    5
15    2.0    8.0
3              All-Bran   K    C        70        4    1
260    9.0    7.0
4 All-Bran_with_Extra_Fiber   K    C        50        4    0
140  14.0    8.0
5              Almond_Delight   R    C        110        2    2
200    1.0  14.0
6 Apple_Cinnamon_Cheerios   G    C        110        2    2
180    1.5  10.5
      sugars potass vitamins shelf weight cups   rating
1         6     280        25     3      1 0.33 68.40297
2         8     135         0     3      1 1.00 33.98368
3         5     320        25     3      1 0.33 59.42551
4         0     330        25     3      1 0.50 93.70491
5         8      -1        25     3      1 0.75 34.38484

```

```
6      10      70      25      1      1 0.75 29.50954
```

▼ Code

```
cereal_tbl |> head(7) # using pipe
```

```
# A tibble: 7 × 16
  name          mfr type calor...1 protein fat sodium fiber
carbo sugars potass
  <chr>         <chr> <chr>   <dbl>   <dbl> <dbl> <dbl> <dbl>
<dbl> <dbl> <dbl>
1 100%_Bran    N     C       70      4     1    130   10
5           6    280
2 100%_Natur... Q     C      120      3     5     15    2
8           8    135
3 All-Bran     K     C       70      4     1    260    9
7           5    320
4 All-Bran_w... K     C       50      4     0    140   14
8           0    330
5 Almond_Del... R     C      110      2     2    200    1
14          8    -1
6 Apple_Cinn... G     C      110      2     2    180   1.5
10.5        10    70
7 Apple_Jacks K     C      110      2     0    125    1
11         14    30
# ... with 5 more variables: vitamins <dbl>, shelf <dbl>, weight
<dbl>,
# cups <dbl>, rating <dbl>, and abbreviated variable name 1
calories
```

▼ Code

```
cereal_tbl |> class() # tells us this is a tibble
```

```
[1] "spec_tbl_df" "tbl_df"      "tbl"         "data.frame"
```

- b. What are the column names of the cereal data frame? How many rows are there? (`dim` and `nrow`)

Solution

▼ Code

```
cereal_tbl |> colnames()
```

```
[1] "name"      "mfr"       "type"      "calories" "protein"
```

```
"fat"
[7] "sodium"  "fiber"    "carbo"    "sugars"   "potass"
"vitamins"
[13] "shelf"    "weight"   "cups"     "rating"
```

▼ Code

```
cereal_tbl |> dim()
```

```
[1] 77 16
```

▼ Code

```
cereal_tbl |> nrow()
```

```
[1] 77
```

▼ Code

```
dim(cereal)
```

```
[1] 77 16
```

▼ Code

```
nrow(cereal)
```

```
[1] 77
```

- c. Extract the `calories` column using the `$` operator and using the `[[` operator.

Solution

▼ Code

```
### Some 'tidy' ways
Cal <- cereal_tbl %>% select(calories) #tibble with one column
Cal <- cereal_tbl %>% pull(calories) #pull out the column

### Base R
AlternativeCal<-cereal[["calories"]]
acal <- cereal$calories
identical(AlternativeCal, acal)
```

```
[1] TRUE
```

▼ Code

```
class(cereal["calories"])
```

```
[1] "data.frame"
```

▼ Code

```
class(cereal[["calories"]])
```

```
[1] "integer"
```

d. Extract rows 1 to 10 from the `cereal` data frame.

Solution

▼ Code

```
### Base R
cereal[1:10,]
```

				name	mfr	type	calories	protein	fat
sodium	fiber	carbo							
1				100%_Bran	N	C	70	4	1
130	10.0	5.0							
2				100%_Natural_Bran	Q	C	120	3	5
15	2.0	8.0							
3				All-Bran	K	C	70	4	1
260	9.0	7.0							
4	All-Bran_with_Extra_Fiber				K	C	50	4	0
140	14.0	8.0							
5				Almond_Delight	R	C	110	2	2
200	1.0	14.0							
6	Apple_Cinnamon_Cheerios				G	C	110	2	2
180	1.5	10.5							
7				Apple_Jacks	K	C	110	2	0
125	1.0	11.0							
8				Basic_4	G	C	130	3	2
210	2.0	18.0							
9				Bran_Chex	R	C	90	2	1
200	4.0	15.0							
10				Bran_Flakes	P	C	90	3	0
210	5.0	13.0							
	sugars	potass	vitamins	shelf	weight	cups	rating		

1	6	280	25	3	1.00	0.33	68.40297
2	8	135	0	3	1.00	1.00	33.98368
3	5	320	25	3	1.00	0.33	59.42551
4	0	330	25	3	1.00	0.50	93.70491
5	8	-1	25	3	1.00	0.75	34.38484
6	10	70	25	1	1.00	0.75	29.50954
7	14	30	25	2	1.00	1.00	33.17409
8	8	100	25	3	1.33	0.75	37.03856
9	6	125	25	1	1.00	0.67	49.12025
10	5	190	25	3	1.00	0.67	53.31381

▼ Code

```
### Tidyverse
cereal_tbl %>% slice(1:10)
```

```
# A tibble: 10 × 16
  name          mfr type calor...1 protein fat sodium fiber
carbo sugars potass
  <chr>         <chr> <chr>   <dbl>   <dbl> <dbl>  <dbl> <dbl>
<dbl> <dbl> <dbl>
1 100%_Bran    N     C       70      4     1    130   10
5          6    280
2 100%_Natu... Q     C      120      3     5     15    2
8          8    135
3 All-Bran    K     C       70      4     1    260    9
7          5    320
4 All-Bran_... K     C       50      4     0    140   14
8          0    330
5 Almond_De... R     C      110      2     2    200    1
14         8    -1
6 Apple_Cin... G     C      110      2     2    180   1.5
10.5        10    70
7 Apple_Jac... K     C      110      2     0    125    1
11        14    30
8 Basic_4     G     C      130      3     2    210    2
18         8    100
9 Bran_Chex   R     C       90      2     1    200    4
15         6    125
10 Bran_Flak... P     C       90      3     0    210    5
13         5    190
# ... with 5 more variables: vitamins <dbl>, shelf <dbl>, weight
<dbl>,
# cups <dbl>, rating <dbl>, and abbreviated variable name 1
calories
```

e. Make a new data frame called **Kelloggs** that only contains rows that

belongs to manufacturer, Kelloggs (when `mfr` takes the value "K").

Solution

▼ Code

```
### Base R
Kelloggs <- subset(cereal, mfr == "K")
head(Kelloggs)
```

		name	mfr	type	calories	protein	fat
sodium	fiber	carbo					
3		All-Bran	K	C	70	4	1
260	9						
4		All-Bran_with_Extra_Fiber	K	C	50	4	0
140	14						
7		Apple_Jacks	K	C	110	2	0
125	1						
17		Corn_Flakes	K	C	100	2	0
290	1						
18		Corn_Pops	K	C	110	1	0
90	1						
20		Cracklin'_Oat_Bran	K	C	110	3	3
140	4						
sugars	potass	vitamins	shelf	weight	cups	rating	
3	5	320	25	3	1 0.33	59.42551	
4	0	330	25	3	1 0.50	93.70491	
7	14	30	25	2	1 1.00	33.17409	
17	2	35	25	1	1 1.00	45.86332	
18	12	20	25	2	1 1.00	35.78279	
20	7	160	25	3	1 0.50	40.44877	

▼ Code

```
Kelloggs.2 <- cereal[cereal$mfr == "K", -2] #removes 2nd c
head(Kelloggs.2)
```

		name	type	calories	protein	fat	sodium
fiber	carbo						
3		All-Bran	C	70	4	1	260
9	7						
4		All-Bran_with_Extra_Fiber	C	50	4	0	140
14	8						
7		Apple_Jacks	C	110	2	0	125
1	11						
17		Corn_Flakes	C	100	2	0	290
1	21						

18		Corn_Pops	C	110	1	0	90	
1	13							
20		Cracklin'_Oat_Bran	C	110	3	3	140	
4	10							
		sugars	potass	vitamins	shelf	weight	cups	rating
3	5	320	25	3	1	0.33	59.42551	
4	0	330	25	3	1	0.50	93.70491	
7	14	30	25	2	1	1.00	33.17409	
17	2	35	25	1	1	1.00	45.86332	
18	12	20	25	2	1	1.00	35.78279	
20	7	160	25	3	1	0.50	40.44877	

▼ Code

```
### Base R splitting
cereal.splitted <- split(cereal, cereal$mfr) #list of 7 it
Kelloggs.3 <- cereal.splitted["K"]
Kelloggs.4 <- cereal.splitted[["K"]]
identical(Kelloggs, Kelloggs.4)
```

[1] TRUE

▼ Code

```
### Tidyverse way
kelloggs_tbl <- cereal_tbl %>% filter(mfr == "K") #tidy wa
kelloggs_tbl <- cereal_tbl %>% filter(mfr == "K") %>% sele
```

Factors

- Load the `Cereal` data again with the `read.csv` command again. This time, use the optional argument, `stringsAsFactors = TRUE`.
- The `mfr` and `type` columns are now factors. Check that this is true.

Solution

▼ Code

```
cereal <- read_csv("Cereal.csv")
```

Rows: 77 Columns: 16
— Column specification

Delimiter: ","
chr (3): name, mfr, type
dbl (13): calories, protein, fat, sodium, fiber, carbo,

sugars, potass, vita...

i Use ``spec()`` to retrieve the full column specification for this data.

i Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

▼ Code

```
cereal.with.factors <- read.csv("Cereal.csv", stringsAsFactors = FALSE)
cereal.with.factors
```

	protein	fat	sodium	name	mfr	type	calories
1				100%_Bran	N	C	70
4	1	130					
2				100%_Natural_Bran	Q	C	120
3	5	15					
3				All-Bran	K	C	70
4	1	260					
4				All-Bran_with_Extra_Fiber	K	C	50
4	0	140					
5				Almond_Delight	R	C	110
2	2	200					
6				Apple_Cinnamon_Cheerios	G	C	110
2	2	180					
7				Apple_Jacks	K	C	110
2	0	125					
8				Basic_4	G	C	130
3	2	210					
9				Bran_Chex	R	C	90
2	1	200					
10				Bran_Flakes	P	C	90
3	0	210					
11				Cap'n'Crunch	Q	C	120
1	2	220					
12				Cheerios	G	C	110
6	2	290					
13				Cinnamon_Toast_Crunch	G	C	120
1	3	210					
14				Clusters	G	C	110
3	2	140					
15				Cocoa_Puffs	G	C	110
1	1	180					
16				Corn_Chex	R	C	110
2	0	280					
17				Corn_Flakes	K	C	100

2	0	290				
18			Corn_Pops	K	C	110
1	0	90				
19			Count_Chocula	G	C	110
1	1	180				
20			Cracklin'_Oat_Bran	K	C	110
3	3	140				
21			Cream_of_Wheat_(Quick)	N	H	100
3	0	80				
22			Crispix	K	C	110
2	0	220				
23			Crispy_Wheat_&_Raisins	G	C	100
2	1	140				
24			Double_Chex	R	C	100
2	0	190				
25			Froot_Loops	K	C	110
2	1	125				
26			Frosted_Flakes	K	C	110
1	0	200				
27			Frosted_Mini-Wheats	K	C	100
3	0	0				
28	Fruit_&_Fibre_Dates,_Walnuts,_and_Oats			P	C	120
3	2	160				
29			Fruitful_Bran	K	C	120
3	0	240				
30			Fruity_Pebbles	P	C	110
1	1	135				
31			Golden_Crisp	P	C	100
2	0	45				
32			Golden_Grahams	G	C	110
1	1	280				
33			Grape_Nuts_Flakes	P	C	100
3	1	140				
34			Grape-Nuts	P	C	110
3	0	170				
35			Great_Grains_Pecan	P	C	120
3	3	75				
36			Honey_Graham_Ohs	Q	C	120
1	2	220				
37			Honey_Nut_Cheerios	G	C	110
3	1	250				
38			Honey-comb	P	C	110
1	0	180				
39			Just_Right_Crunchy__Nuggets	K	C	110
2	1	170				
40			Just_Right_Fruit_&_Nut	K	C	140
3	1	170				
41			Kix	G	C	110

2	1	260				
42			Life	Q	C	100
4	2	150				
43			Lucky_Charms	G	C	110
2	1	180				
44			Maypo	A	H	100
4	1	0				
45			Muesli_Raisins,_Dates,_&_Almonds	R	C	150
4	3	95				
46			Muesli_Raisins,_Peaches,_&_Pecans	R	C	150
4	3	150				
47			Mueslix_Crispy_Blend	K	C	160
3	2	150				
48			Multi-Grain_Cheerios	G	C	100
2	1	220				
49			Nut&Honey_Crunch	K	C	120
2	1	190				
50			Nutri-Grain_Almond-Raisin	K	C	140
3	2	220				
51			Nutri-grain_Wheat	K	C	90
3	0	170				
52			Oatmeal_Raisin_Crisp	G	C	130
3	2	170				
53			Post_Nat._Raisin_Bran	P	C	120
3	1	200				
54			Product_19	K	C	100
3	0	320				
55			Puffed_Rice	Q	C	50
1	0	0				
56			Puffed_Wheat	Q	C	50
2	0	0				
57			Quaker_Oat_Squares	Q	C	100
4	1	135				
58			Quaker_Oatmeal	Q	H	100
5	2	0				
59			Raisin_Bran	K	C	120
3	1	210				
60			Raisin_Nut_Bran	G	C	100
3	2	140				
61			Raisin_Squares	K	C	90
2	0	0				
62			Rice_Chex	R	C	110
1	0	240				
63			Rice_Krispies	K	C	110
2	0	290				
64			Shredded_Wheat	N	C	80
2	0	0				
65			Shredded_Wheat_'n'Bran	N	C	90

3	0	0							
66			Shredded_Wheat_spoon_size	N	C			90	
3	0	0							
67			Smacks	K	C			110	
2	1	70							
68			Special_K	K	C			110	
6	0	230							
69			Strawberry_Fruit_Wheats	N	C			90	
2	0	15							
70			Total_Corn_Flakes	G	C			110	
2	1	200							
71			Total_Raisin_Bran	G	C			140	
3	1	190							
72			Total_Whole_Grain	G	C			100	
3	1	200							
73			Triples	G	C			110	
2	1	250							
74			Trix	G	C			110	
1	1	140							
75			Wheat_Chex	R	C			100	
3	1	230							
76			Wheaties	G	C			100	
3	1	200							
77			Wheaties_Honey_Gold	G	C			110	
2	1	200							
	fiber	carbo	sugars	potass	vitamins	shelf	weight	cups	
	rating								
1	10.0	5.0	6	280	25	3	1.00	0.33	
68.40297									
2	2.0	8.0	8	135	0	3	1.00	1.00	
33.98368									
3	9.0	7.0	5	320	25	3	1.00	0.33	
59.42551									
4	14.0	8.0	0	330	25	3	1.00	0.50	
93.70491									
5	1.0	14.0	8	-1	25	3	1.00	0.75	
34.38484									
6	1.5	10.5	10	70	25	1	1.00	0.75	
29.50954									
7	1.0	11.0	14	30	25	2	1.00	1.00	
33.17409									
8	2.0	18.0	8	100	25	3	1.33	0.75	
37.03856									
9	4.0	15.0	6	125	25	1	1.00	0.67	
49.12025									
10	5.0	13.0	5	190	25	3	1.00	0.67	
53.31381									
11	0.0	12.0	12	35	25	2	1.00	0.75	

18.04285							
12	2.0	17.0	1	105	25	1	1.00 1.25
50.76500							
13	0.0	13.0	9	45	25	2	1.00 0.75
19.82357							
14	2.0	13.0	7	105	25	3	1.00 0.50
40.40021							
15	0.0	12.0	13	55	25	2	1.00 1.00
22.73645							
16	0.0	22.0	3	25	25	1	1.00 1.00
41.44502							
17	1.0	21.0	2	35	25	1	1.00 1.00
45.86332							
18	1.0	13.0	12	20	25	2	1.00 1.00
35.78279							
19	0.0	12.0	13	65	25	2	1.00 1.00
22.39651							
20	4.0	10.0	7	160	25	3	1.00 0.50
40.44877							
21	1.0	21.0	0	-1	0	2	1.00 1.00
64.53382							
22	1.0	21.0	3	30	25	3	1.00 1.00
46.89564							
23	2.0	11.0	10	120	25	3	1.00 0.75
36.17620							
24	1.0	18.0	5	80	25	3	1.00 0.75
44.33086							
25	1.0	11.0	13	30	25	2	1.00 1.00
32.20758							
26	1.0	14.0	11	25	25	1	1.00 0.75
31.43597							
27	3.0	14.0	7	100	25	2	1.00 0.80
58.34514							
28	5.0	12.0	10	200	25	3	1.25 0.67
40.91705							
29	5.0	14.0	12	190	25	3	1.33 0.67
41.01549							
30	0.0	13.0	12	25	25	2	1.00 0.75
28.02576							
31	0.0	11.0	15	40	25	1	1.00 0.88
35.25244							
32	0.0	15.0	9	45	25	2	1.00 0.75
23.80404							
33	3.0	15.0	5	85	25	3	1.00 0.88
52.07690							
34	3.0	17.0	3	90	25	3	1.00 0.25
53.37101							
35	3.0	13.0	4	100	25	3	1.00 0.33

45.81172								
36	1.0	12.0	11	45	25	2	1.00	1.00
21.87129								
37	1.5	11.5	10	90	25	1	1.00	0.75
31.07222								
38	0.0	14.0	11	35	25	1	1.00	1.33
28.74241								
39	1.0	17.0	6	60	100	3	1.00	1.00
36.52368								
40	2.0	20.0	9	95	100	3	1.30	0.75
36.47151								
41	0.0	21.0	3	40	25	2	1.00	1.50
39.24111								
42	2.0	12.0	6	95	25	2	1.00	0.67
45.32807								
43	0.0	12.0	12	55	25	2	1.00	1.00
26.73451								
44	0.0	16.0	3	95	25	2	1.00	1.00
54.85092								
45	3.0	16.0	11	170	25	3	1.00	1.00
37.13686								
46	3.0	16.0	11	170	25	3	1.00	1.00
34.13976								
47	3.0	17.0	13	160	25	3	1.50	0.67
30.31335								
48	2.0	15.0	6	90	25	1	1.00	1.00
40.10596								
49	0.0	15.0	9	40	25	2	1.00	0.67
29.92429								
50	3.0	21.0	7	130	25	3	1.33	0.67
40.69232								
51	3.0	18.0	2	90	25	3	1.00	1.00
59.64284								
52	1.5	13.5	10	120	25	3	1.25	0.50
30.45084								
53	6.0	11.0	14	260	25	3	1.33	0.67
37.84059								
54	1.0	20.0	3	45	100	3	1.00	1.00
41.50354								
55	0.0	13.0	0	15	0	3	0.50	1.00
60.75611								
56	1.0	10.0	0	50	0	3	0.50	1.00
63.00565								
57	2.0	14.0	6	110	25	3	1.00	0.50
49.51187								
58	2.7	-1.0	-1	110	0	1	1.00	0.67
50.82839								
59	5.0	14.0	12	240	25	2	1.33	0.75

```

39.25920
60  2.5 10.5      8   140      25    3   1.00 0.50
39.70340
61  2.0 15.0      6   110      25    3   1.00 0.50
55.33314
62  0.0 23.0      2    30      25    1   1.00 1.13
41.99893
63  0.0 22.0      3    35      25    1   1.00 1.00
40.56016
64  3.0 16.0      0    95       0    1   0.83 1.00
68.23588
65  4.0 19.0      0   140       0    1   1.00 0.67
74.47295
66  3.0 20.0      0   120       0    1   1.00 0.67
72.80179
67  1.0  9.0     15    40      25    2   1.00 0.75
31.23005
68  1.0 16.0      3    55      25    1   1.00 1.00
53.13132
69  3.0 15.0      5    90      25    2   1.00 1.00
59.36399
70  0.0 21.0      3    35     100    3   1.00 1.00
38.83975
71  4.0 15.0     14   230     100    3   1.50 1.00
28.59278
72  3.0 16.0      3   110     100    3   1.00 1.00
46.65884
73  0.0 21.0      3    60      25    3   1.00 0.75
39.10617
74  0.0 13.0     12    25      25    2   1.00 1.00
27.75330
75  3.0 17.0      3   115      25    1   1.00 0.67
49.78744
76  3.0 17.0      3   110      25    1   1.00 1.00
51.59219
77  1.0 16.0      8    60      25    1   1.00 0.75
36.18756

```

▼ Code

```
levels(cereal.with.factors$mfr)
```

```
[1] "A" "G" "K" "N" "P" "Q" "R"
```

▼ Code

```
class(cereal.with.factors$mfr)
```

```
[1] "factor"
```

▼ Code

```
class(cereal$mfr)
```

```
[1] "character"
```

▼ Code

```
# or
class(cereal.with.factors$carbo)
```

```
[1] "numeric"
```

▼ Code

```
class(cereal$carbo)
```

```
[1] "numeric"
```

▼ Code

```
# or
str(cereal.with.factors) #only characters become factors
```

```
'data.frame': 77 obs. of 16 variables:
 $ name      : Factor w/ 77 levels
"100%_Bran","100%_Natural_Bran",...: 1 2 3 4 5 6 7 8 9 10 ...
 $ mfr       : Factor w/ 7 levels "A","G","K","N",...: 4 6 3 3 7
2 3 2 7 5 ...
 $ type      : Factor w/ 2 levels "C","H": 1 1 1 1 1 1 1 1 1 1
...
 $ calories: int 70 120 70 50 110 110 110 130 90 90 ...
 $ protein : int 4 3 4 4 2 2 2 3 2 3 ...
 $ fat      : int 1 5 1 0 2 2 0 2 1 0 ...
 $ sodium   : int 130 15 260 140 200 180 125 210 200 210 ...
 $ fiber     : num 10 2 9 14 1 1.5 1 2 4 5 ...
 $ carbo     : num 5 8 7 8 14 10.5 11 18 15 13 ...
 $ sugars    : int 6 8 5 0 8 10 14 8 6 5 ...
 $ potass    : int 280 135 320 330 -1 70 30 100 125 190 ...
 $ vitamins: int 25 0 25 25 25 25 25 25 25 25 ...
 $ shelf     : int 3 3 3 3 3 1 2 3 1 3 ...
 $ weight    : num 1 1 1 1 1 1 1 1.33 1 1 ...
 $ cups      : num 0.33 1 0.33 0.5 0.75 0.75 1 0.75 0.67 0.67
...
...
```



```
$ rating : num 68.4 34 59.4 93.7 34.4 ...
```

▼ Code

```
str(cereal)
```

```
spc_tbl_ [77 × 16] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
 $ name      : chr [1:77] "100%_Bran" "100%_Natural_Bran" "All-
Bran" "All-Bran_with_Extra_Fiber" ...
 $ mfr       : chr [1:77] "N" "Q" "K" "K" ...
 $ type      : chr [1:77] "C" "C" "C" "C" ...
 $ calories: num [1:77] 70 120 70 50 110 110 110 130 90 90 ...
 $ protein  : num [1:77] 4 3 4 4 2 2 2 3 2 3 ...
 $ fat       : num [1:77] 1 5 1 0 2 2 0 2 1 0 ...
 $ sodium   : num [1:77] 130 15 260 140 200 180 125 210 200 210
...
 $ fiber     : num [1:77] 10 2 9 14 1 1.5 1 2 4 5 ...
 $ carbo     : num [1:77] 5 8 7 8 14 10.5 11 18 15 13 ...
 $ sugars    : num [1:77] 6 8 5 0 8 10 14 8 6 5 ...
 $ potass    : num [1:77] 280 135 320 330 -1 70 30 100 125 190
...
 $ vitamins: num [1:77] 25 0 25 25 25 25 25 25 25 25 ...
 $ shelf     : num [1:77] 3 3 3 3 3 1 2 3 1 3 ...
 $ weight    : num [1:77] 1 1 1 1 1 1 1 1.33 1 1 ...
 $ cups      : num [1:77] 0.33 1 0.33 0.5 0.75 0.75 1 0.75 0.67
0.67 ...
 $ rating    : num [1:77] 68.4 34 59.4 93.7 34.4 ...
- attr(*, "spec")=
.. cols(
..   name = col_character(),
..   mfr = col_character(),
..   type = col_character(),
..   calories = col_double(),
..   protein = col_double(),
..   fat = col_double(),
..   sodium = col_double(),
..   fiber = col_double(),
..   carbo = col_double(),
..   sugars = col_double(),
..   potass = col_double(),
..   vitamins = col_double(),
..   shelf = col_double(),
..   weight = col_double(),
..   cups = col_double(),
..   rating = col_double()
.. )
- attr(*, "problems")=<externalptr>
```

- b. How many levels are there in `mfr` and `type`? (use the functions `levels` or `nlevels`)

Solution

▼ Code

```
levels(cereal.with.factors$mfr)
```

```
[1] "A" "G" "K" "N" "P" "Q" "R"
```

▼ Code

```
# or  
nlevels(cereal.with.factors$mfr)
```

```
[1] 7
```

▼ Code

```
# or  
str(cereal.with.factors$mfr)
```

```
Factor w/ 7 levels "A","G","K","N",...: 4 6 3 3 7 2 3 2 7 5  
...
```

▼ Code

```
# class() typeof()
```

Vectors

- a. Extract the `calories` into a new vector called `cereal.calories`.

Solution

▼ Code

```
cereal.calories <- cereal$calories  
cereal.calories <- cereal[["calories"]]  
cereal_calories <- cereal_tbl %>% pull(calories)
```

- b. How many elements are there in `cereal.calories`? (`length`)

Solution

▼ Code

```
length(cereal.calories)
```

```
[1] 77
```

▼ Code

```
cereal_calories %>% length()
```

```
[1] 77
```

c. Extract the 5th to the 10th element from `cereal.calories`.

Solution

▼ Code

```
# cereal.calories[5:10] # most code works this way  
cereal_calories[5:10]
```

```
[1] 110 110 110 130 90 90
```

d. Add one more element to `cereal.calories` using `c()`.

Solution

▼ Code

```
cereal_calories <- c(cereal_calories, 1.0) #c for concater  
length(cereal.calories)
```

```
[1] 77
```

Matrix

a. Can you force the cereal data frame to be a Matrix?
(`as.matrix(cereal)`). Check that the elements have been forced into the character type.

Solution

▼ Code

```
cereal.matrix <- as.matrix(cereal)  
str(cereal.matrix)
```

```
chr [1:77, 1:16] "100%_Bran" "100%_Natural_Bran" "All-Bran"
...
- attr(*, "dimnames")=List of 2
..$ : NULL
..$ : chr [1:16] "name" "mfr" "type" "calories" ...
```

- b. Now do this again, but this time leave out the `mfr`, `name` and `type` columns. Check that the elements are now numeric.

Solution

▼ Code

```
cereal.removed <- cereal[, -(1:3)]
cereal.removed
```

```
# A tibble: 77 × 13
  calories protein fat sodium fiber carbo sugars potass
vitamins shelf weight
      <dbl>   <dbl> <dbl>   <dbl> <dbl> <dbl>   <dbl>   <dbl>
<dbl> <dbl>   <dbl>
1       70       4     1    130  10     5       6    280
25      3      1
2      120       3     5     15   2     8       8    135
0      3      1
3       70       4     1    260   9     7       5    320
25      3      1
4       50       4     0    140  14     8       0    330
25      3      1
5      110       2     2    200   1    14       8     -1
25      3      1
6      110       2     2    180  1.5  10.5     10     70
25      1      1
7      110       2     0    125   1    11      14     30
25      2      1
8      130       3     2    210   2    18       8    100
25      3    1.33
9       90       2     1    200   4    15       6    125
25      1      1
10      90       3     0    210   5    13       5    190
25      3      1
# ... with 67 more rows, and 2 more variables: cups <dbl>,
rating <dbl>
```

▼ Code

```
cereal.numeric.matrix <- as.matrix(cereal.removed)
str(cereal.numeric.matrix)
```

```
num [1:77, 1:13] 70 120 70 50 110 110 110 130 90 90 ...
- attr(*, "dimnames")=List of 2
..$ : NULL
..$ : chr [1:13] "calories" "protein" "fat" "sodium" ...
```

1.3 Numerical summary

Summary

- Use the `summary` function to extract the median, 1st quartile and 3rd quartile data from the `sodium` column.

Solution

▼ Code

```
summary(cereal$sodium)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0	130.0	180.0	159.7	210.0	320.0

Basic statistics

- Find the max, min, standard deviation and mean of the `sodium` (`max()`, `min()`, `sd()`, `mean()`)

Solution

▼ Code

```
cereal_tbl %>% pull(sodium) %>% summary()
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0	130.0	180.0	159.7	210.0	320.0

▼ Code

```
sodium<-cereal$sodium
max(sodium)
```

```
[1] 320
```

▼ Code

```
min(cereal$sodium)
```

```
[1] 0
```

▼ Code

```
sd(cereal$sodium)
```

```
[1] 83.8323
```

▼ Code

```
mean(cereal$sodium)
```

```
[1] 159.6753
```

▼ Code

```
summary(sodium)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0	130.0	180.0	159.7	210.0	320.0

c. Find the mean `sodium` of each `mfr`.

Solution

▼ Code

```
### Can be done by repeated subsetting, this is tedious
kelloggs.cereals <- subset(cereal, mfr == "K")
mean(kelloggs.cereals$sodium)
```

```
[1] 174.7826
```

▼ Code

```
### Can use a formula and the aggregate function

mean.sodiums <- aggregate(sodium ~ mfr, data = cereal, FUN = mean)
### Can split vector (or data.frame if you wanted) by
### another vector. In this case, split by species.
```

```
split.sodium <- split(cereal$sodium, cereal$mfr)
### Apply a function over a list and return a list (_l_app
lapply(split.sodium, mean)
```

```
$A
[1] 0
```

```
$G
[1] 200.4545
```

```
$K
[1] 174.7826
```

```
$N
[1] 37.5
```

```
$P
[1] 146.1111
```

```
$Q
[1] 92.5
```

```
$R
[1] 198.125
```

▼ Code

```
### Apply a function over a list and return a _simplified
sapply(split.sodium, mean)
```

	A	G	K	N	P	Q	R
	0.0000	200.4545	174.7826	37.5000	146.1111	92.5000	198.1250

▼ Code

```
### Also could use by and tapply, vapply, for the interest
```

```
cereal_tbl %>%
  select(sodium, mfr) %>%
  group_by(mfr) %>%
  summarise(mean_sodium = mean(sodium))
```

```
# A tibble: 7 × 2
  mfr    mean_sodium
<chr>      <dbl>
1 A              0
2 G          200.
```

3	K	175.
4	N	37.5
5	P	146.
6	Q	92.5
7	R	198.

1.4 Graphical summary

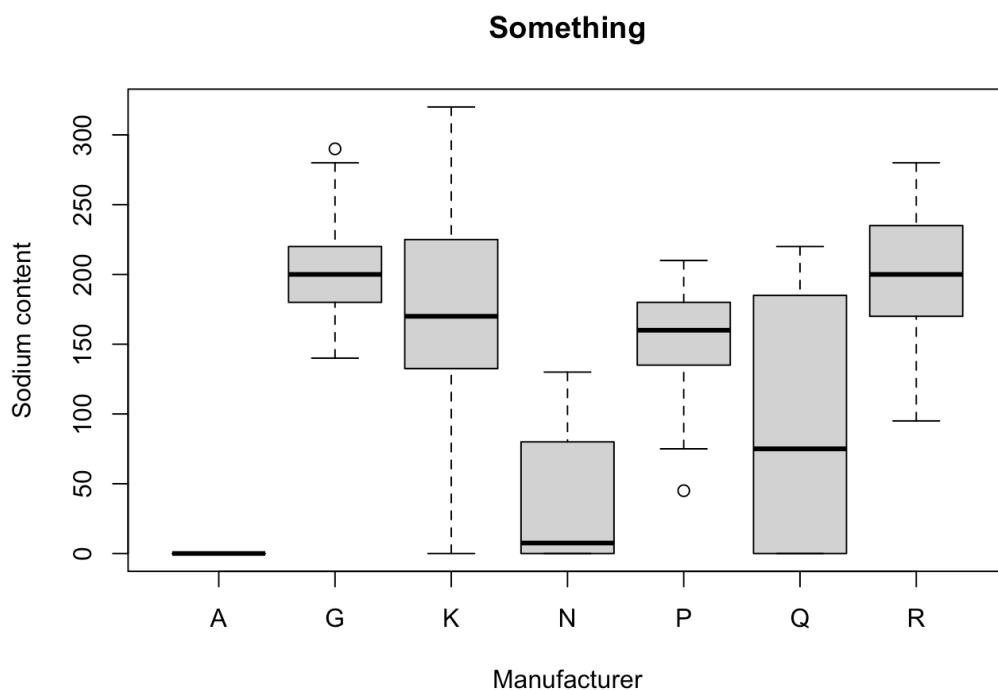
Boxplot

- a. Make a boxplot of the `sodium` against `mfr` using `boxplot()`.

Solution

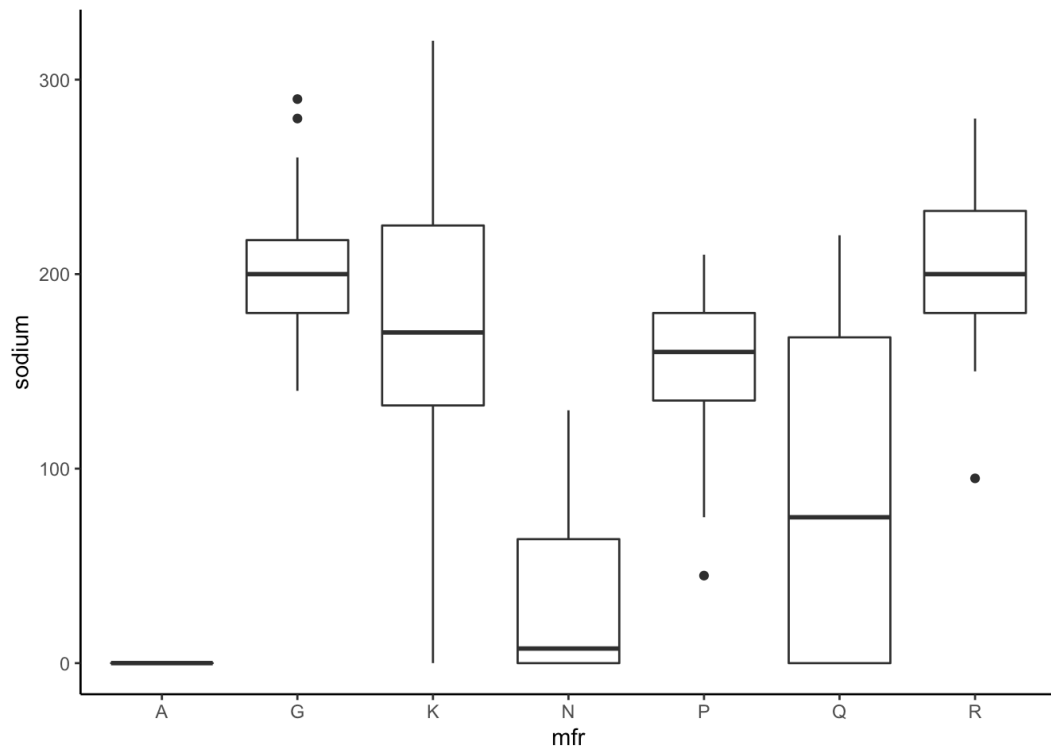
▼ Code

```
boxplot(sodium ~ mfr, data = cereal,
        xlab = "Manufacturer", ylab = "Sodium content", ma
```



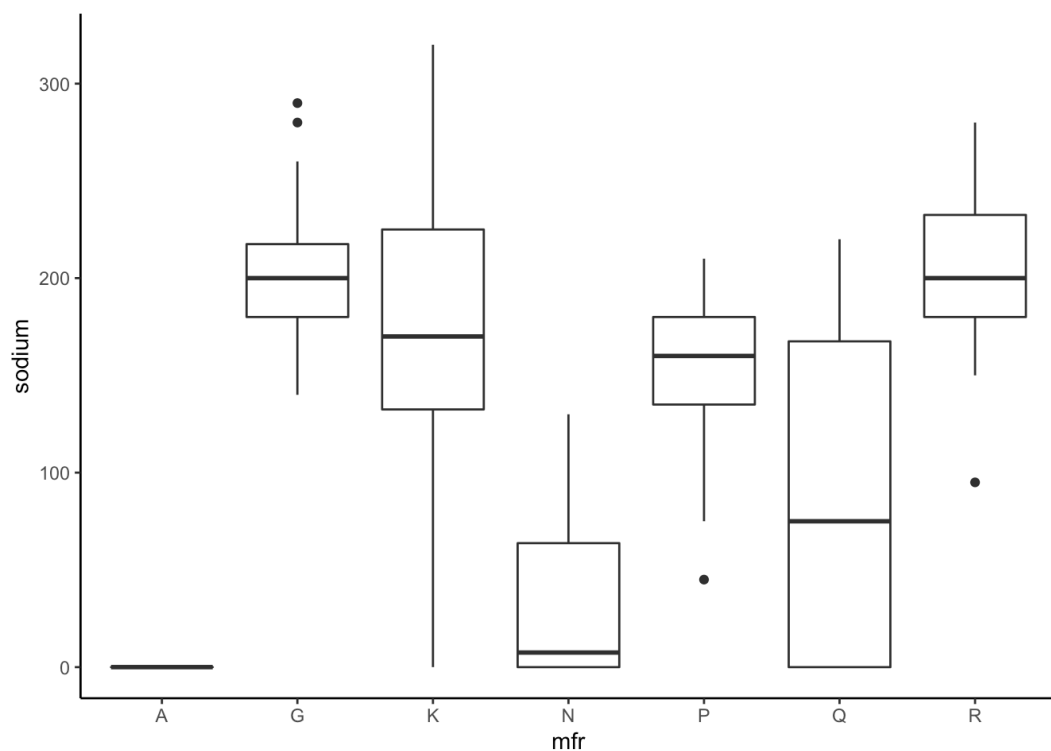
▼ Code

```
ggplot(cereal_tbl, aes(x = mfr, y = sodium)) + geom_boxplot
```

▼ Code

```
cereal_tbl %>% ggplot(aes(x = mfr, y = sodium)) + geom_boxplot()
```



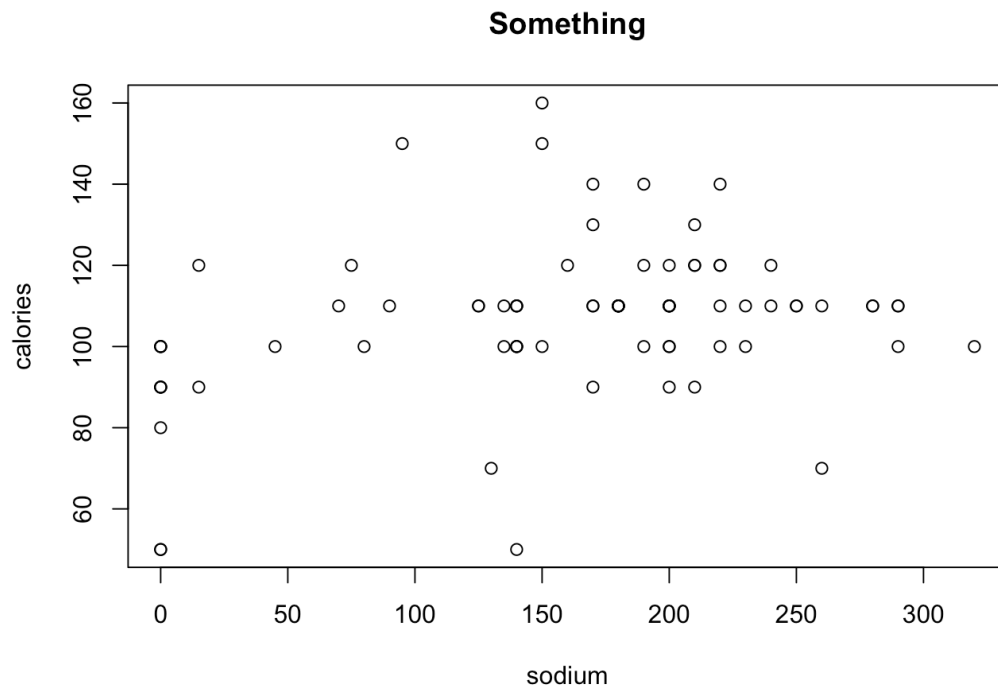
Scatterplot

b. Plot `calories` against `sodium` using `plot()`.

Solution

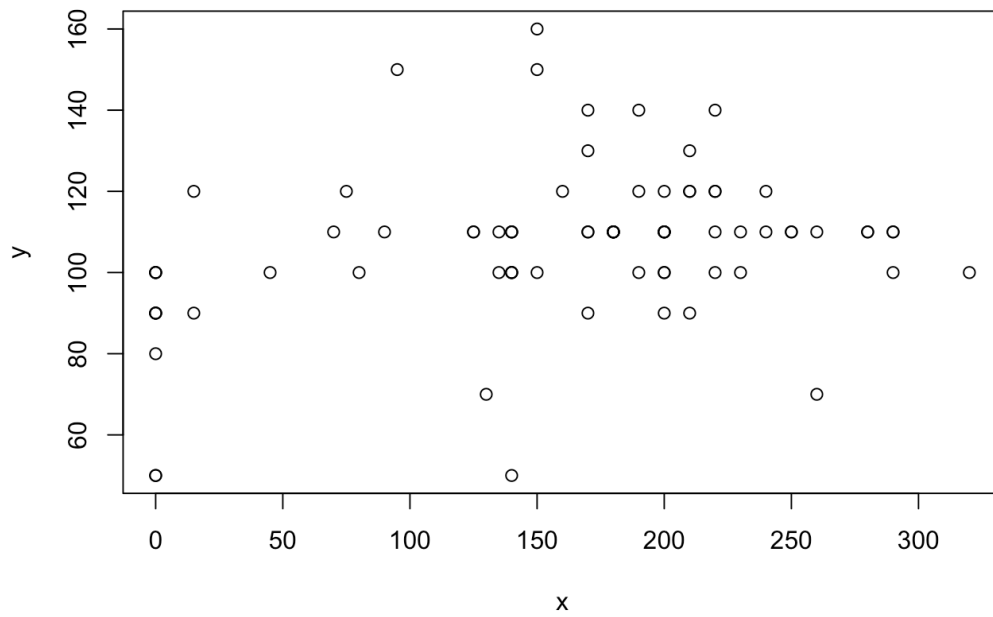
▼ Code

```
### Using formula  
plot(calories ~ sodium, data = cereal, main = "Something")
```



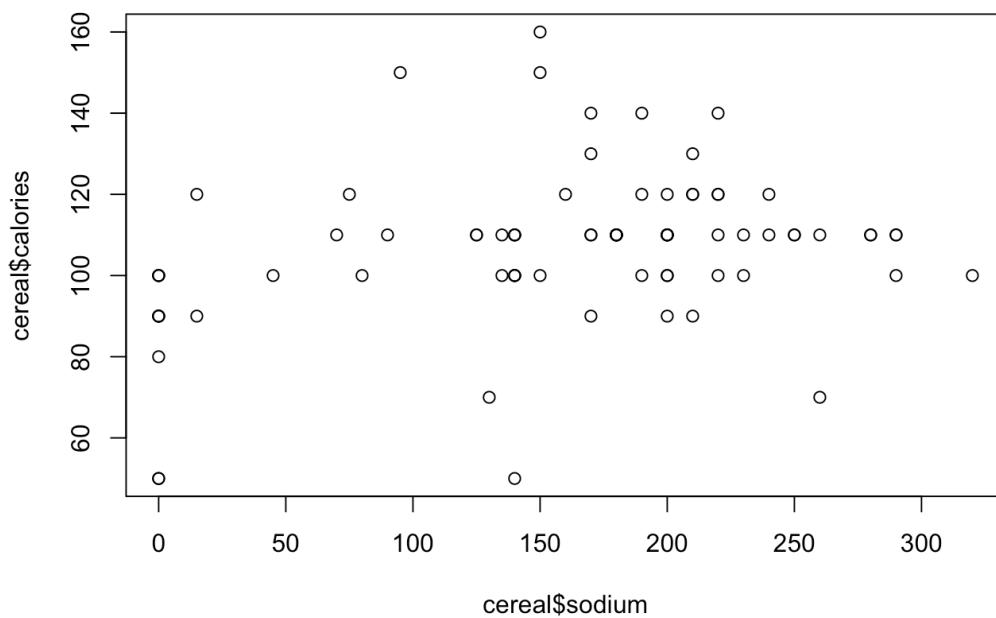
▼ Code

```
### Another way, define x and y  
x <- cereal$sodium  
y <- cereal$calories  
plot(x, y)
```



▼ Code

```
plot(cereal$sodium, cereal$calories)
```

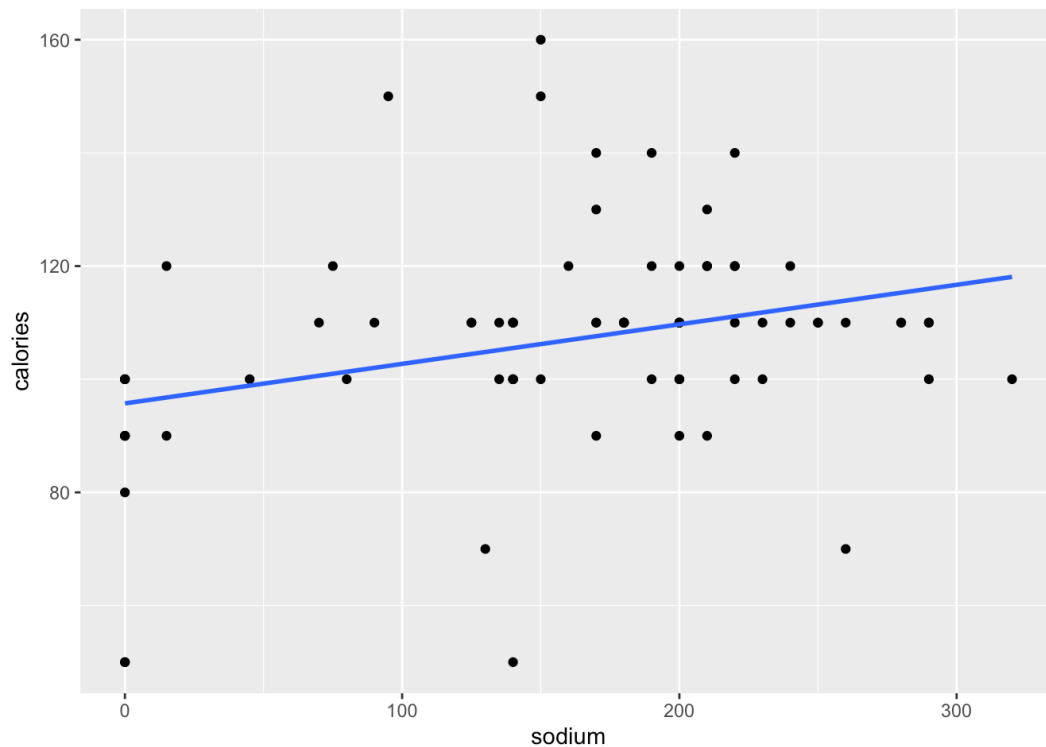


▼ Code

```
### Alternatively, use with to help R find the vectors  
###with(cereal, plot(sodium, calories))
```

```
ggplot(cereal_tbl, aes(x = sodium, y = calories)) +
  geom_point() + geom_smooth(method = "lm", se = FALSE)
```

`geom_smooth()` using formula 'y ~ x'



1.5 Write Data to File

- b. Write data frame with only the Kellogg's observations to a file called `kelloggs.csv`. Use the `write.csv` command.

Solution

▼ Code

```
write.csv(kelloggs_tbl, file = "kelloggs.csv")
head(cereal)
```

```
# A tibble: 6 × 16
  name      mfr  type calor...1 protein  fat sodium fiber
carbo sugars potass
  <chr>      <chr> <chr>   <dbl>   <dbl> <dbl> <dbl> <dbl>
<dbl> <dbl>   <dbl>
1 100%_Bran  N      C       70      4      1    130    10
5          6 280
2 100%_Natur... Q      C      120      3      5     15     2
8          8 135
```

```

3 All-Bran      K      C      70      4      1      260      9
7              5      320
4 All-Bran_w... K      C      50      4      0      140      14
8              0      330
5 Almond_Del... R      C      110     2      2      200      1
14             8      -1
6 Apple_Cinn... G      C      110     2      2      180      1.5
10.5           10      70
# ... with 5 more variables: vitamins <dbl>, shelf <dbl>, weight
<dbl>,
# cups <dbl>, rating <dbl>, and abbreviated variable name 1
calories

```

2 Melbourne house prices regression model

In this section we will examine the dataset describing Melbourne house prices. This dataset was downloaded from [Kaggle](#) and the data was released under the CC BY-NC-SA 4.0 license. For this lab, we will focus on three subbrubs - Brunswick, Craigieburn and Hawthorn and examine what variables or factors are associated with the housing price.

2.1 Load the data

Load the Melbourne house price dataset from Canvas.

Solution

▼ Code

```

melb.dat <- read.csv("Melbourne_housing_FULL.csv")
melbdata <- read_csv("Melbourne_housing_FULL.csv")

```

Rows: 34857 Columns: 21

— Column specification

Delimiter: ","

chr (8): Suburb, Address, Type, Method, SellerG, Date, CouncilArea, Regionname

dbl (13): Rooms, Price, Distance, Postcode, Bedroom2, Bathroom, Car, Landsiz...

i Use ``spec()`` to retrieve the full column specification for this data.

i Specify the column types or set ``show_col_types = FALSE`` to

quiet this message.

2.2 Initial data analysis

We will need to subset the data to only look at 3 suburbs - Brunswick, Craigieburn and Hawthorn. Similar to lab 1, start the data analysis by generating some quantitative and graphical summaries. For example, determine the average price in each of these three suburbs. Explore more summaries of the data.

Solution

▼ Code

```
### Base R
melb.data.sub <- subset(melbdata, Suburb == "Hawthorn" | S
melb.data.sub2 <- subset(melbdata, Suburb %in% c("Hawthorn
identical(melb.data.sub, melb.data.sub2)
```

[1] TRUE

▼ Code

```
split.data <- split(melb.data.sub[["Price"]], melb.data.sub
suburb.means <- vapply(split.data, mean, numeric(1L), na.rm
suburb.medians <- vapply(split.data, median, numeric(1L),
```

▼ Code

```
### Tidyverse way
melbdata.sub <- melbdata %>%
  filter(Suburb %in% c("Hawthorn", "Brunswick", "Craigieburn
  mutate(Suburb = factor(Suburb, levels = c("Craigieburn", "Brunswick", "Hawthorn"))

melbdata %>%
  filter(Suburb %in% c("Hawthorn", "Brunswick", "Craigieburn
  group_by(Suburb) %>%
  summarise(Mean_Price = mean(Price, na.rm = TRUE), Median_Price = median(Price, na.rm = TRUE))
```

A tibble: 3 × 3

Suburb	Mean_Price	Median_price
<chr>	<dbl>	<dbl>
1 Brunswick	977989.	950000
2 Craigieburn	566173.	562500
3 Hawthorn	1238074.	750500

For the following questions, use the subsetting data for the Suburbs of Brunswick, Craigieburn and Hawthorn.

2.3 Finding association I

To examine the association between house prices and a single variable, start by constructing a simple linear regression using only **BuildingArea** as a predictor. Use an appropriate statistic to justify the goodness of fit of the prediction and create a graphical output to enable you to assess your model fit.

Note: you might consider other variables too.

Solution Consider a scatter plot of **Price** against **BuildingArea** and overlay the prediction from the linear regression model.

▼ Code

```
### FORM THE LINEAR REGRESSION MODEL
lm1 <- lm(data = melbdata.sub, Price/1000 ~ BuildingArea)

### Inspect coefficients
coef(lm1)
```

```
(Intercept) BuildingArea
518.192115      3.800746
```

▼ Code

```
lm1 |> coef()#get coefficients
```

```
(Intercept) BuildingArea
518.192115      3.800746
```

▼ Code

```
lm1 |> fitted() |> head() #fitted values
```

```
      2      3      4      9     11     12
928.6727 871.6615 1312.5480 1609.0062 769.0413 670.2220
```

▼ Code

```
lm1 |> resid() |> head() #residuals ie errors, check mean
```

2 3 4 9 11

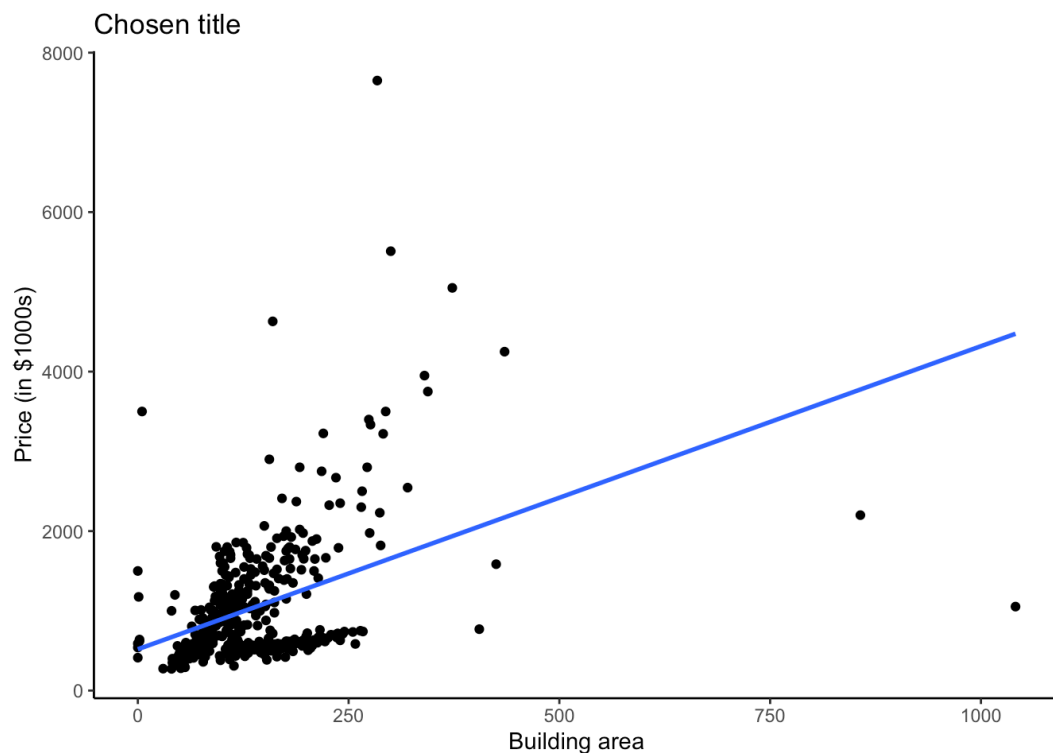
12

97.32732 930.83851 187.45198 620.99380 -359.04135

-397.72195

▼ Code

```
ggplot(melbdata.sub |> select(BuildingArea, Price) |> drop
  aes(x = BuildingArea, y = Price/1000) +
  geom_point() + geom_smooth(formula = y ~ x, method = '
  theme_classic() + labs(x = "Building area", y = "Price
```



▼ Code

```
### Base R way
summary(lm1)
```

Call:

```
lm(formula = Price/1000 ~ BuildingArea, data = melbdata.sub)
```

Residuals:

Min	1Q	Median	3Q	Max
-3421.8	-463.9	-148.0	259.1	6052.4

Coefficients:

Estimate	Std. Error	t value	Pr(> t)
----------	------------	---------	----------


```
(Intercept) 518.1921    66.5956    7.781 5.74e-14 ***
BuildingArea  3.8007     0.4082    9.311 < 2e-16 ***
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 730 on 416 degrees of freedom

(709 observations deleted due to missingness)

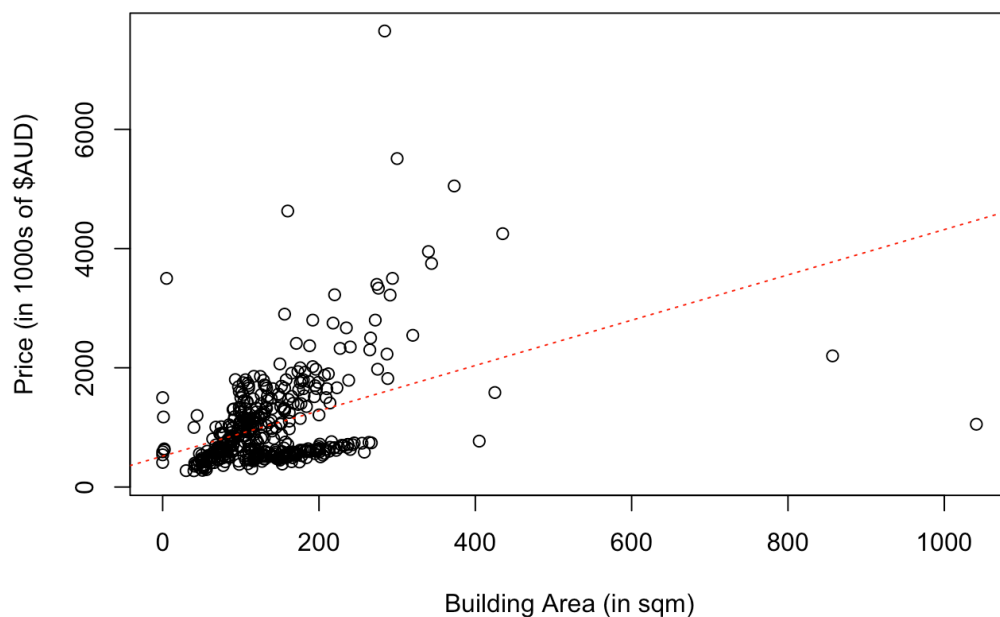
Multiple R-squared: 0.1725, Adjusted R-squared: 0.1705

F-statistic: 86.69 on 1 and 416 DF, p-value: < 2.2e-16

▼ Code

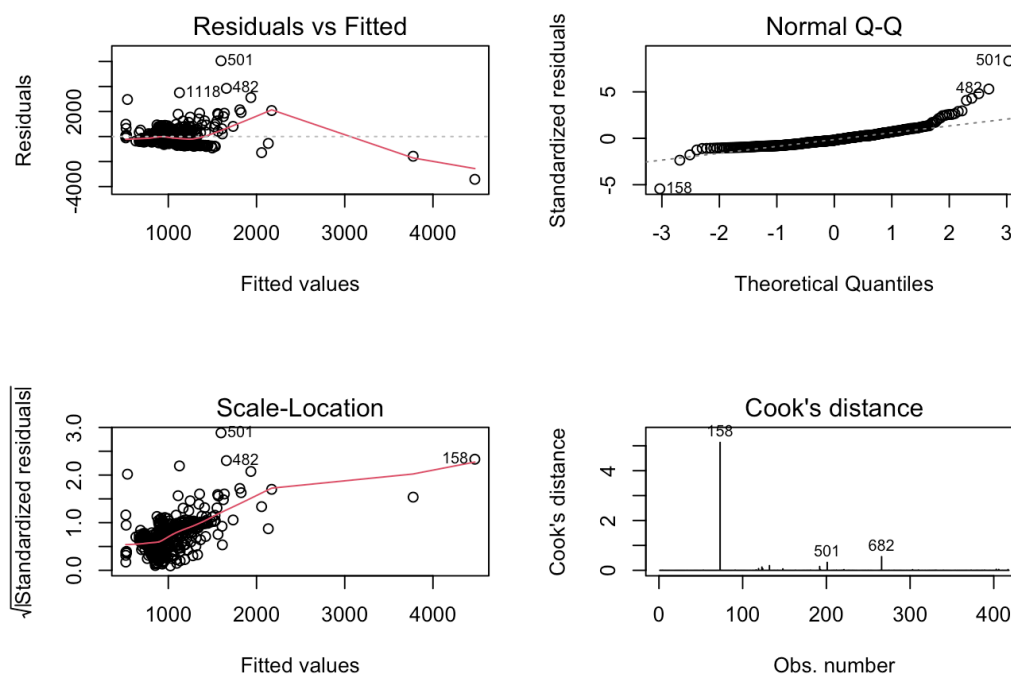
```
plot(Price/1000 ~ BuildingArea, data = melbdata.sub,
     main = "House prices in Brunswick, Craigieburn and Ha
     xlab = "Building Area (in sqm)", ylab = "Price (in 10
     abline(lm1, col = "red", lty = "dotted")
```

House prices in Brunswick, Craigieburn and Hawthorn



▼ Code

```
par(mfrow = c(2, 2))
plot(lm1, which = 1:4)
```



▼ Code

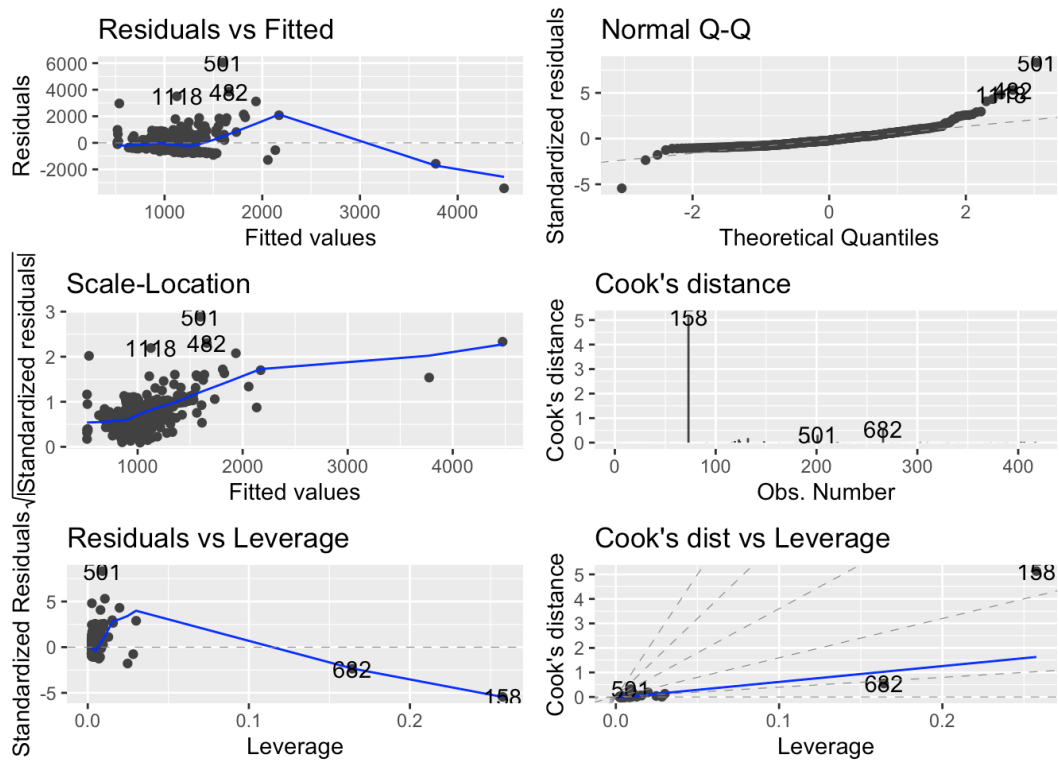
```
r2 <- round(summary(lm1)$r.squared, 4)
r2
```

```
[1] 0.1725
```

There is 17.25% of the variation in Price explained by the linear regression on Building Area.

▼ Code

```
### Tidyverse way
library(ggfortify)
autoplot(lm1, which = 1:6, nrow = 3, ncol = 2)
```



▼ Code

```
### LOTS OF NICE CONVENIENT CODE
### lm1 |> coefficients() #get coefficients
### lm1 |> fitted() #fitted values
### lm1 |> residuals() %>% mean()#residuals ie errors, check for bias

library(broom)
lm1 |> augment() #full table of fitted values, cooks distance, etc
```

```
# A tibble: 418 × 8
  .rownames `Price/1000` BuildingArea .fitted .hat .sigma
  .cooks.d .std...1
  <chr>          <dbl>          <dbl> <dbl> <dbl> <dbl>
<dbl> <dbl>
1 2              1026          108  929. 0.00267 731.
0.0000238 0.134
2 3              1802.           93  872. 0.00302 729.
0.00247 1.28
3 4              1500          209 1313. 0.00398 731.
0.000132 0.257
4 9              2230          287 1609. 0.00936 730.
0.00345 0.855
5 11             410           66  769. 0.00400 731.
0.000488 -0.493
6 12             272.           40  670. 0.00538 731.
0.000807 -0.546
7 13             680          100  898. 0.00284 731.
```

```
0.000128 -0.299
 8 16          400          61    750. 0.00423    731.
0.000491 -0.481
 9 17          950          96    883. 0.00294    731.
0.0000124 0.0918
10 20          860          97    887. 0.00291    731.
0.00000198 -0.0369
# ... with 408 more rows, and abbreviated variable name 1
.std.resid
```

▼ Code

```
lm1 |> glance() #key values eg R squared, can pull out
```

```
# A tibble: 1 × 12
  r.squared adj.r.squa...1 sigma stati...2 p.value    df logLik
AIC      BIC devia...3
  <dbl>      <dbl> <dbl>    <dbl>    <dbl> <dbl> <dbl>
<dbl> <dbl>    <dbl>
1    0.172      0.170  730.     86.7 7.41e-19    1 -3348.
6702. 6714.  2.22e8
# ... with 2 more variables: df.residual <int>, nobs <int>, and
abbreviated
# variable names 1adj.r.squared, 2statistic, 3deviance
```

▼ Code

```
lm1 |> tidy() #conveniently puts summary into tibble format
```

```
# A tibble: 2 × 5
  term          estimate std.error statistic p.value
<chr>          <dbl>    <dbl>    <dbl>    <dbl>
1 (Intercept)    518.      66.6      7.78 5.74e-14
2 BuildingArea     3.80      0.408     9.31 7.41e-19
```

▼ Code

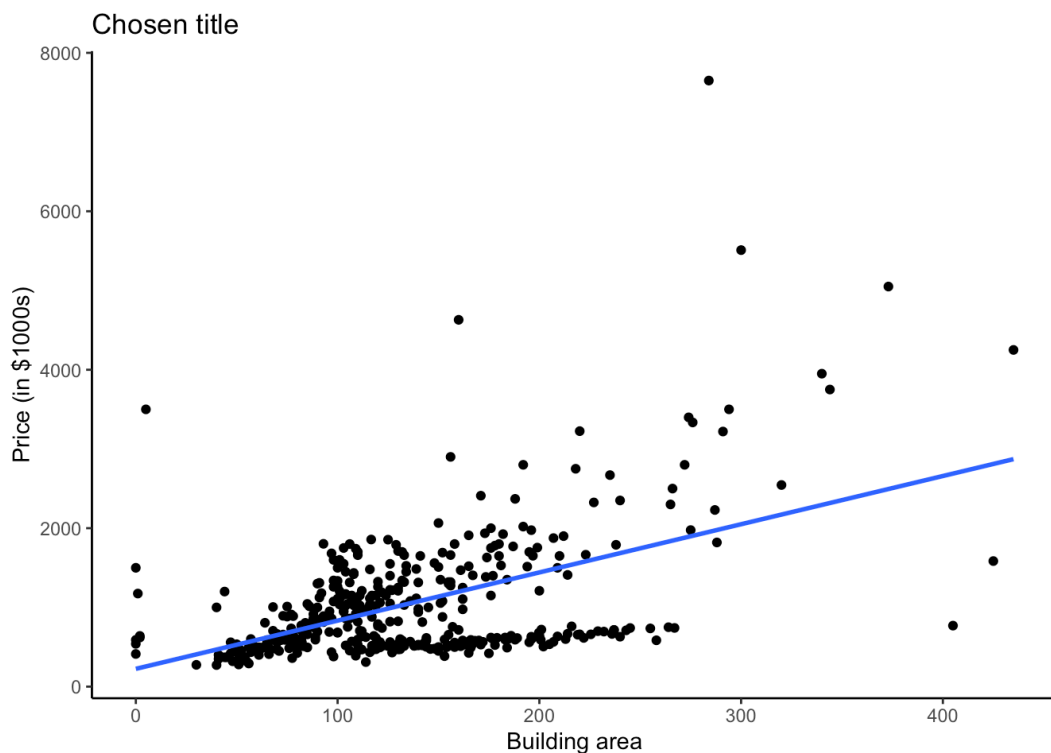
```
r2 <- lm1 |> glance() |> pull(r.squared)

melbdata.sub_out <- melbdata %>%
  filter(Suburb %in% c("Hawthorn", "Brunswick", "Craigieburn"))
  mutate(Suburb = factor(Suburb, levels = c("Craigieburn", "Hawthorn", "Brunswick")))
  slice(-c(158,682)) #REMOVE THE WORST TWO OUTLIERS

lm1_alt <- lm(data = melbdata.sub_out, Price/1000 ~ BuildingArea)

###OUTLIERS HAVE BEEN REMOVED
```

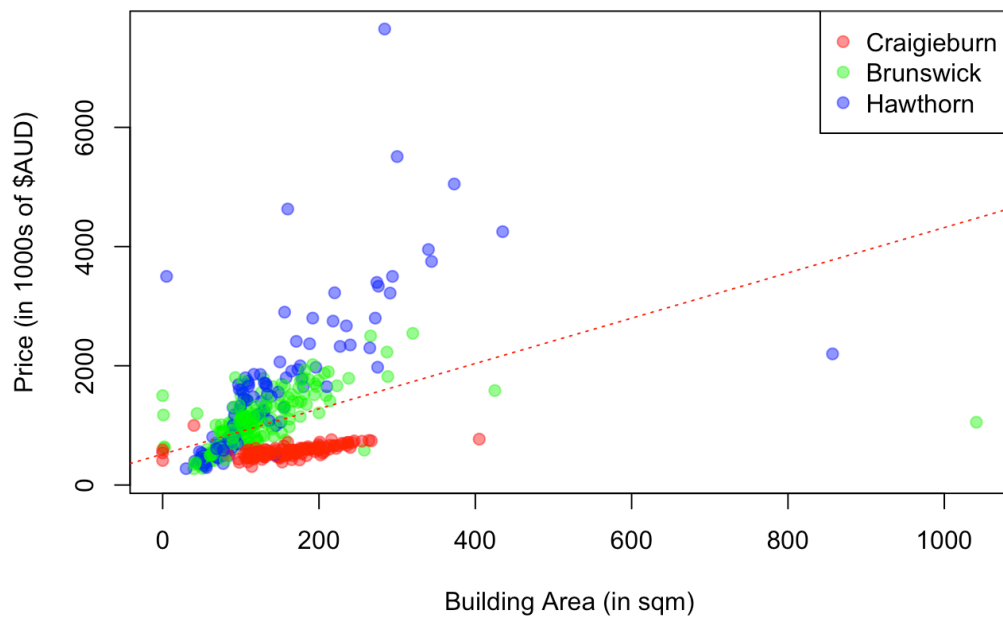
```
ggplot(melbdata.sub_out |> select(BuildingArea, Price) |>
  aes(x = BuildingArea, y = Price/1000) +
  geom_point() + geom_smooth(formula = y ~ x, method = "lm") +
  theme_classic() + labs(x = "Building area", y = "Price")
```



▼ Code

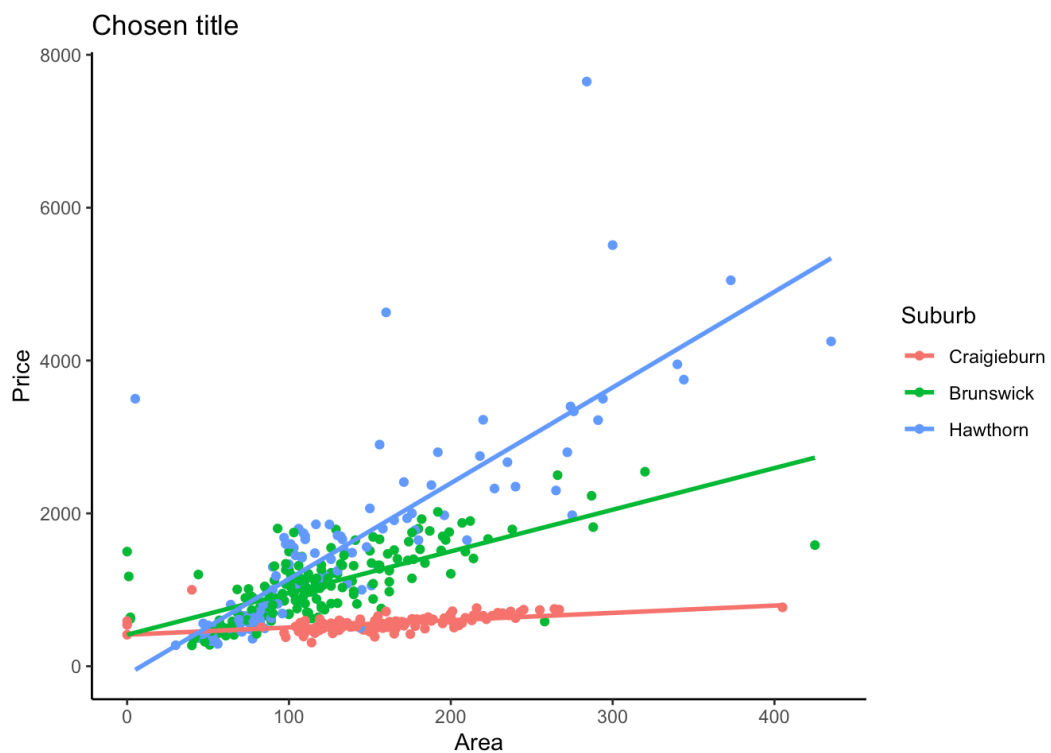
```
# Base R
lmfit1 <- lm(data = melbdata.sub, Price/1000 ~ BuildingArea)
###Get semi-transparent red green and blue
my.colours <- rgb(c(1, 0, 0), c(0, 1, 0), c(0, 0, 1), alpha = 0.5)
plot(Price/1000 ~ BuildingArea, data = melbdata.sub,
     main = "House prices of some suburbs against Building Area",
     xlab = "Building Area (in sqm)", ylab = "Price (in 1000s)",
     col = my.colours[as.integer(melbdata.sub[["Suburb"]])],
     pch = 19)
legend("topright", legend = levels(melbdata.sub[["Suburb"]]),
     col = my.colours, pch = 19)
abline(lmfit1, col = "red", lty = "dotted")
```

House prices of some suburbs against Building Area



▼ Code

```
ggplot(melbdata.sub_out |> select(BuildingArea, Price, Suburb)) +
  aes(x = BuildingArea, y = Price/1000, color = Suburb) +
  geom_point() +
  theme_classic() +
  labs(x = "Area", y = "Price", title = "Chosen title") +
  geom_smooth(formula = y ~ x, method = "lm", se = FALSE)
```



2.4 Finding association II

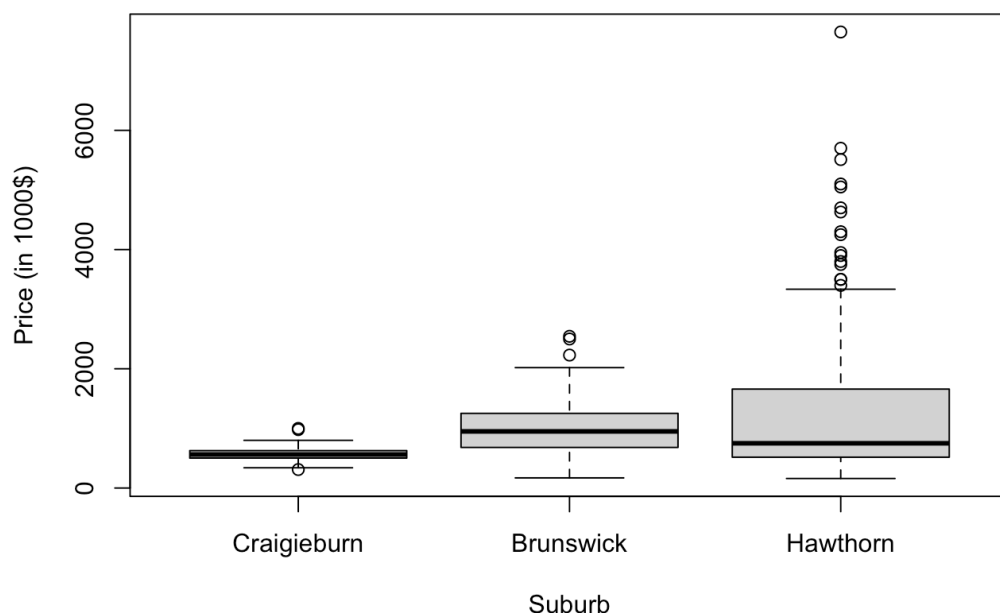
- Variability of house prices are complex and likely to be explained by many different factors. Construct a multiple linear regression here by examining if adding **Suburb** as a predictor will improve the prediction? Notice that **Suburb** is a categorical variable. Briefly describe how to interpret the regression coefficients returned by `lm`.
- There are many other variables in the data, you might consider whether adding the number of car spaces as a predictor improve the prediction model?

Solution

- Model fit below

▼ Code

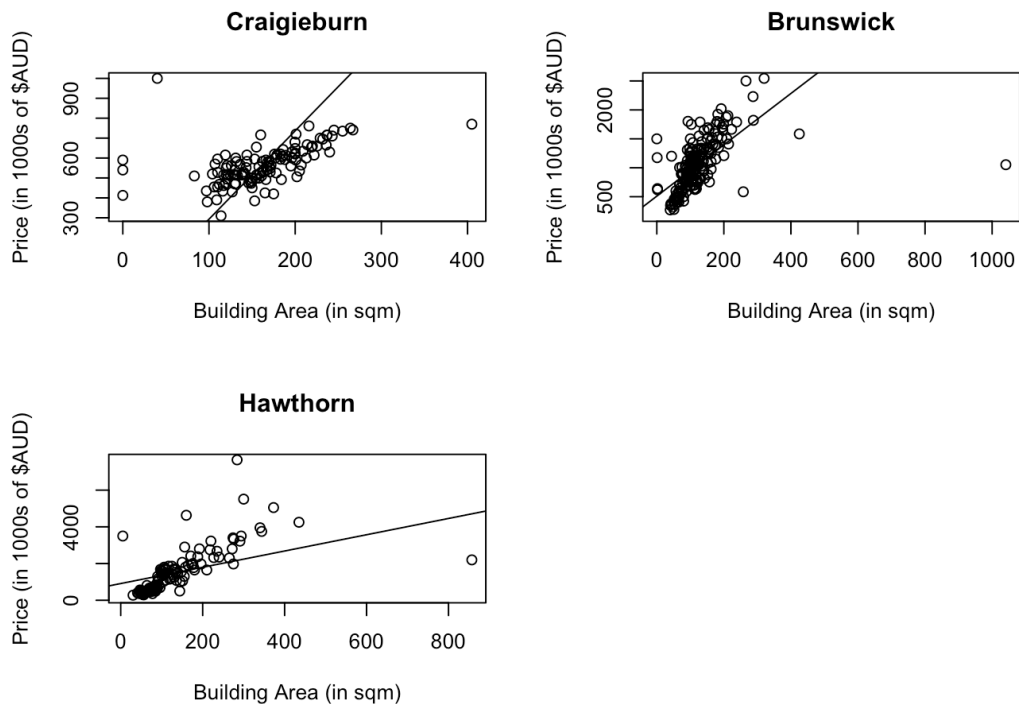
```
# Base R  
boxplot(Price/1000 ~ Suburb, data = melbdata.sub, ylab = 'Price (in 1000$)')
```



▼ Code

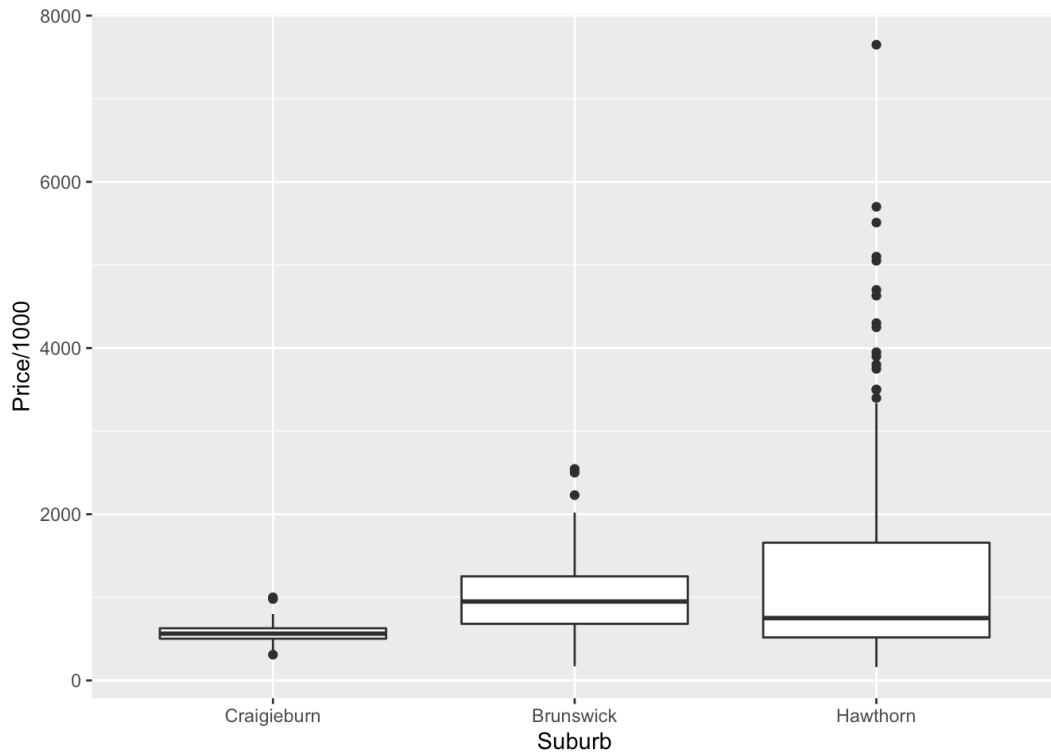
```
lm2 <- lm(Price/1000 ~ BuildingArea + Suburb, data = melb.  
coefs <- lm2 |> coef()
```

```
# Base R
par(mfrow = c(2, 2))
invisible(lapply(levels(melbdata.sub$Suburb), function(x)
  plot(Price/1000 ~ BuildingArea, data = subset(melbdata, Suburb == x),
    xlab = "Building Area (in sqm)", ylab = "Price (in 1000s of $AUD)",
    int <- coefs[1]
    if (any(adjust.ind <- grepl(paste0(x, "$"), names(coef))))
      int <- int + coefs[adjust.ind]
    abline(int, coefs[2])
  )))
```



▼ Code

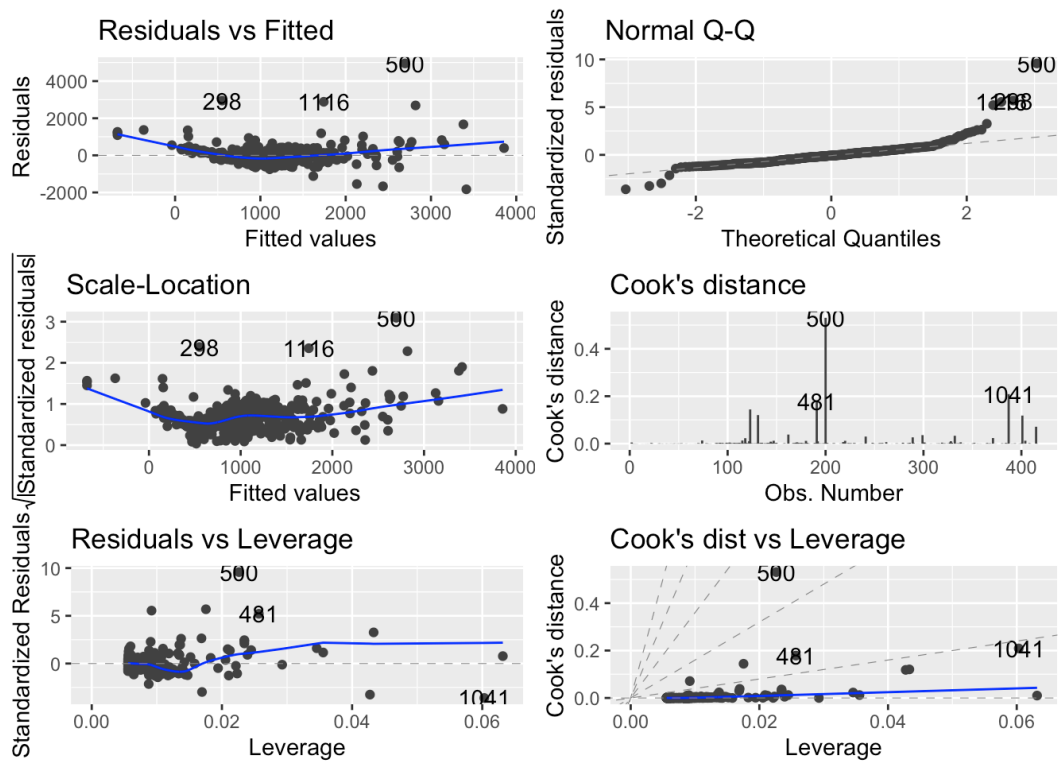
```
# Tidyverse way
ggplot(melbdata.sub |> select(Suburb, Price) |> drop_na())
  aes(x = Suburb, y = Price/1000) + geom_boxplot()
```

▼ Code

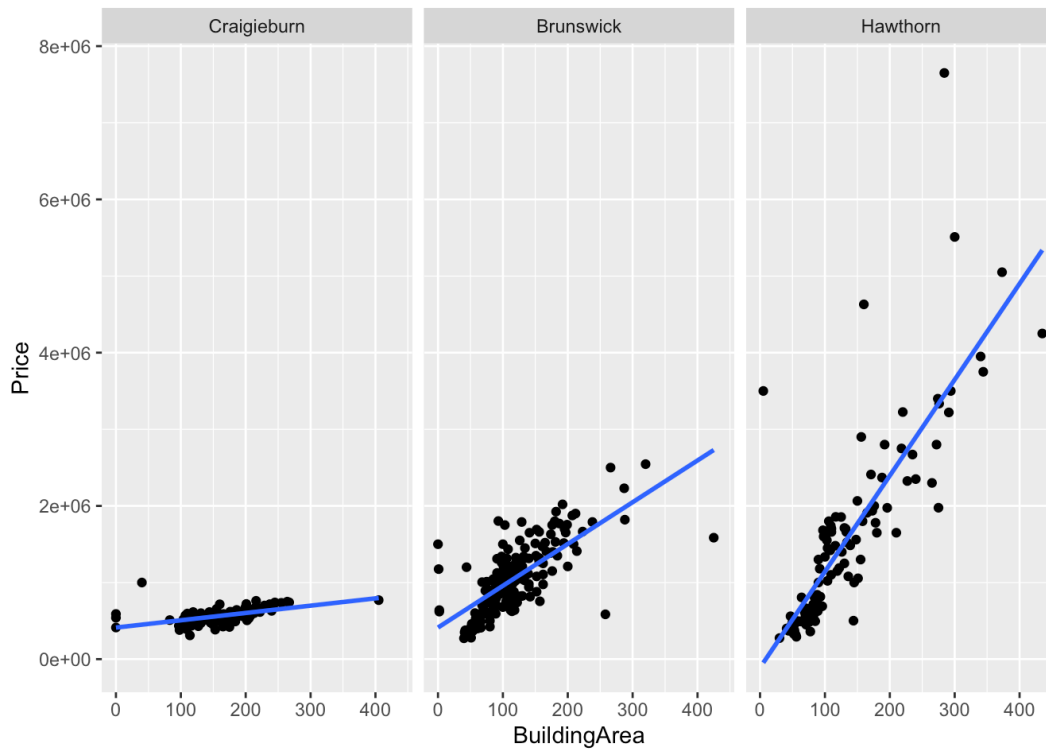
```
lmfit2 <- lm(data = melbdata.sub %>% slice(-c(158,682)), Pr
coefs <- lmfit2$coefficients

autoplot(lmfit2, which = 1:6, nrow = 3, ncol = 2)
```



▼ Code

```
ggplot(melbdata.sub_out |> select(BuildingArea, Price, Suburb)) +
  aes(x = BuildingArea, y = Price) +
  geom_point() +
  geom_smooth(formula = y ~ x, method = "lm", se = FALSE) +
  facet_wrap(~Suburb)
```



▼ Code

```
summary(lmfit2)
```

Call:

```
lm(formula = Price/1000 ~ BuildingArea + Suburb, data =
melbdata.sub %>%
  slice(-c(158, 682)))
```

Residuals:

Min	1Q	Median	3Q	Max
-1829.6	-262.7	-41.5	185.2	4953.7

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-675.6983	80.1195	-8.434	5.71e-16 ***
BuildingArea	7.6864	0.3991	19.260	< 2e-16 ***
SuburbBrunswick	823.6223	63.9521	12.879	< 2e-16 ***
SuburbHawthorn	1189.0488	69.2335	17.174	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 523.2 on 412 degrees of freedom
(709 observations deleted due to missingness)

Multiple R-squared: 0.5768, Adjusted R-squared: 0.5737

F-statistic: 187.2 on 3 and 412 DF, p-value: < 2.2e-16

▼ Code

```
lmfit2 %>% glance()
```

A tibble: 1 × 12

	r.squared	adj.r.squa... ¹	sigma stati... ²	p.value	df	logLik	AIC	BIC	devia... ³
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.577	0.574	523.	187.	1.47e-76	3	-3192.	6395.	6415.

... with 2 more variables: df.residual <int>, nobs <int>, and abbreviated

variable names ¹adj.r.squared, ²statistic, ³deviance

▼ Code

```
r2s <- lmfit2 %>% glance() %>% pull(r.squared)
r2s <- summary(lmfit2)$r.squared
```

One way to highlight that the regression lines for the three suburbs are parallel is to put all three on the same graph, as follows.

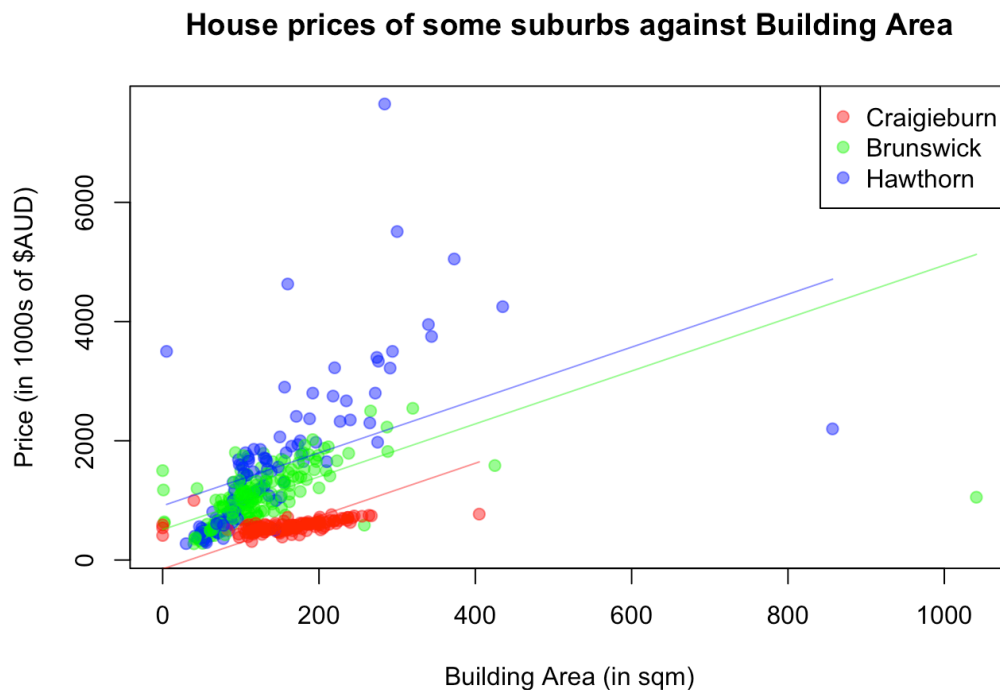
▼ Code

```
# Base R
plot(Price/1000 ~ BuildingArea, data = melbdata.sub,
     main = "House prices of some suburbs against Building Area",
     xlab = "Building Area (in sqm)", ylab = "Price (in 1000s)",
     col = my.colours[as.integer(melbdata.sub[["Suburb"]])],
     pch = 19)
legend("topright", legend = levels(melbdata.sub[["Suburb"]]),
      col = my.colours, pch = 19)
coefs <- coefficients(lmfit2)
names(my.colours) <- levels(melbdata.sub$Suburb)
r2 <- round(summary(lmfit1)$r.squared, 4)
obs.buildingarea.suburb <- subset(melbdata.sub, select = c("BuildingArea", "Price"))
obs.buildingarea.suburb <- na.omit(obs.buildingarea.suburb)
buildingarea.by.suburb <- with(obs.buildingarea.suburb, sp
```

```

buildingarea.by.suburb <- lapply(buildingarea.by.suburb, r
invisible(lapply(levels(melbdata.sub$Suburb), function(x)
  pred.df <- data.frame(BuildingArea = buildingarea.by.s
                        Suburb = x)
  lines(pred.df[["BuildingArea"]], predict(lm2, newdata
    col = my.colours[which(levels(obs.buildingarea.s
}))

```



Embedding model fit in ggplot

`ggplot` includes a really clever trick to support easily constructing line fits (smoothed or linear or ...) however the interactions between the model specification and the plotting specification can be subtle, resulting in graphs that do not match the numerical analysis, which are then at best misleading.

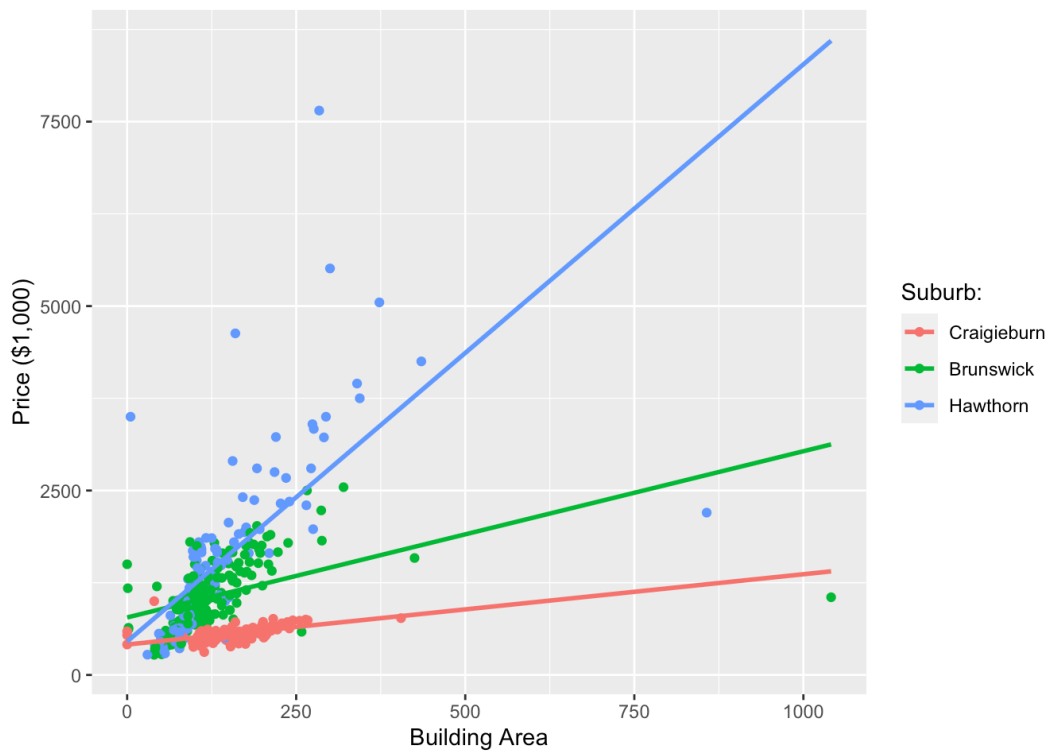
An example of fitting different models.

▼ Code

```

####
library(ggplot2)
ggplot(melbdata.sub, aes(x = BuildingArea, y = Price/1000,
  geom_point(na.rm = TRUE) +
  geom_smooth(formula = "y~x", method = "lm", se = FALSE,
  xlab("Building Area") + ylab("Price ($1,000)") + labs(cc

```



Specifying the colour slightly differently, results in a different model!

▼ Code

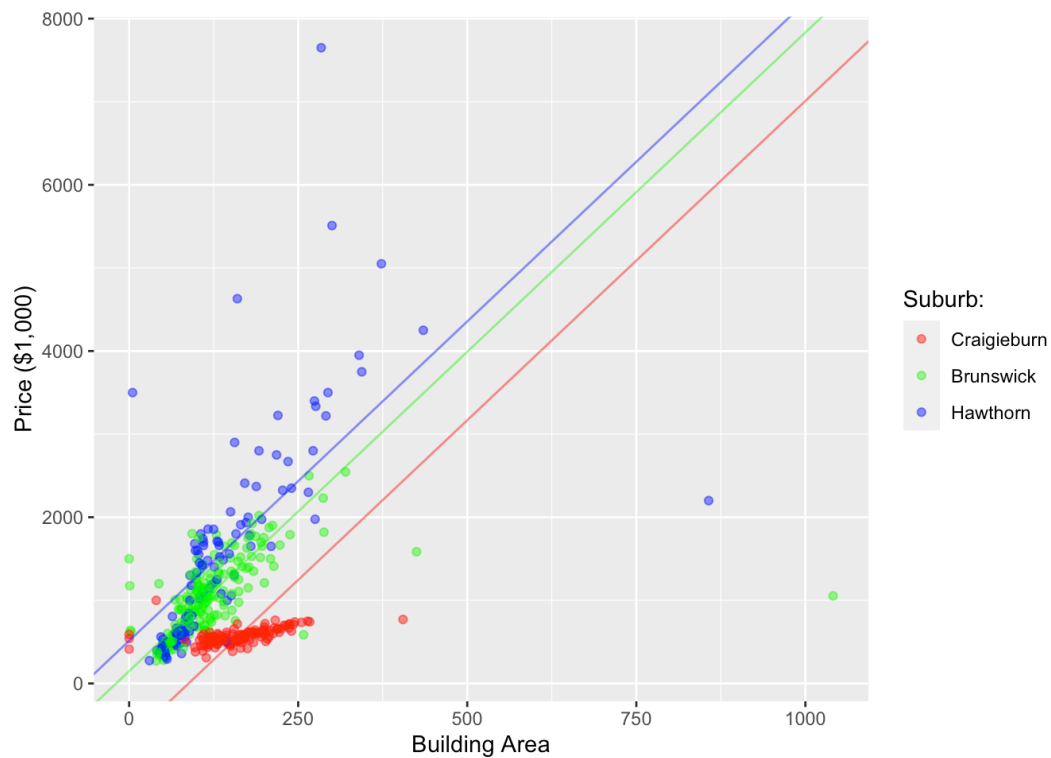
```
ggplot(melbdata.sub, aes(x = BuildingArea, y = Price/1000))  
  geom_point(aes(color=Suburb), na.rm = TRUE) +  
  geom_smooth(formula="y~x", method = "lm", se = FALSE, fu  
  xlab("Building Area") + ylab("Price ($1,000)") + labs(cc
```



A simple way to avoid the problem is to make sure that the model being plotted is the original model used for the numerical analysis:

▼ Code

```
### reuse previous color key
colScale <- scale_colour_manual(name = "Suburb:", values =
ggplot(melbdata.sub, aes(x = BuildingArea, y = Price/1000,
  geom_point(na.rm = TRUE) +
  xlab("Building Area") + ylab("Price ($1,000)") + colScale
  sapply(unique(melbdata.sub$Suburb), function(x) {
    int <- coefs[1]
    if (any(adjust.ind <- grepl(paste0(x, "$"), names(coefs)))
      int <- int + coefs[adjust.ind]
    geom_abline( intercept=int, slope=coefs[2], col=my.col[x])
  })
```



Commentary on Models

The **Suburb** predictor improves the fit of the model by increasing the R^2 from 0.1725 to 0.5767645. However, the adjusted R^2 is a more appropriate goodness of fit measure when there is more than one predictor in the model since adding another predictor will always increase the R^2 . In this case the adjusted R^2 increases by a similar amount suggesting Suburb is a good additional predictor.

Interpreting the **BuildingArea** slope has the interpretation that for each unit increase in square meter of building size, the expected average price would increase by \$4435. Interpreting the categorical predictors needs to be done by intercept adjustment. The first categorical level of **Suburb** (Brunswick) becomes the baseline intercept and the other suburbs are adjusted against the baseline intercept. In this case Craigieburn and Hawthorn have adjustments of -660,000 and 400,000 respectively. This should be interpreted that properties in Craigieburn are \$660,000 cheaper than Brunswick on average (if BuildingArea is held fixed). Hawthorn properties are \$400,000 more expensive than Brunswick. This is consistent with the graphical summary in the boxplot which indicates without adjusting for Building Area, Craigieburn tends to have cheaper houses with low variance while Hawthorn has a large variance in house prices with many very expensive outlying properties.

b. Adding the number of car spaces in to the model and compare the

goodness of fit measures.

▼ Code

```
lmfit4 <- lm(data = melbdata.sub, Price ~ BuildingArea + Suburb + Car, data = melbdata.sub)
summary(lmfit4)
```

Call:

```
lm(formula = Price ~ BuildingArea + Suburb + Car, data = melbdata.sub)
```

Residuals:

Min	1Q	Median	3Q	Max
-3417281	-273352	-59191	252474	5001704

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-447547.1	89421.2	-5.005	8.37e-07 ***
BuildingArea	3761.7	351.3	10.708	< 2e-16 ***
SuburbBrunswick	781107.7	73776.4	10.588	< 2e-16 ***
SuburbHawthorn	1144544.5	78263.7	14.624	< 2e-16 ***
Car	220740.5	34616.8	6.377	4.98e-10 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 588500 on 402 degrees of freedom
(720 observations deleted due to missingness)

Multiple R-squared: 0.477, Adjusted R-squared: 0.4718

F-statistic: 91.66 on 4 and 402 DF, p-value: < 2.2e-16

Adding car spaces seems to improve the prediction model if **BuildingArea** and **Suburb** are already included in the model. The goodness of fit metrics (both raw and adjusted) increase.

2.5 Impact of outliers

Model construction can be affected by unwanted variation and noise such as outliers. For example, houses with very small building areas of 5sqm and lower and larger places over 300 sqm look like outliers. How would you assess the impact of outliers?

Solution

A simple strategy to assess the impact of outliers is to remove the outliers and see if you can improve the prediction model.

▼ Code

```

melbdata.sub.2 <- subset(melbdata.sub, BuildingArea > 10 &
lmfit3 <- lm(data = melbdata.sub.2, Price/1000 ~ BuildingArea)
coefs <- lmfit3$coefficients
par(mfrow = c(2, 2))
invisible(lapply(unique(melbdata.sub$Suburb), function(x)
  plot(Price/1000 ~ BuildingArea, data = subset(melbdata.sub, Suburb == x),
    xlab = "Building Area (in sqm)", ylab = "Price (in thousands)",
    int <- coefs[1]
    if (any(adjust.ind <- grepl(paste0(x, "$"), names(coefs))))
      int <- int + coefs[adjust.ind]
    abline(int, coefs[2])
  )))
summary(lmfit3)

```

Call:

```
lm(formula = Price/1000 ~ BuildingArea + Suburb, data = melbdata.sub.2)
```

Residuals:

Min	1Q	Median	3Q	Max
-1658.0	-260.8	-30.5	192.6	4895.4

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-832.4033	80.4592	-10.35	<2e-16 ***
BuildingArea	8.5448	0.4234	20.18	<2e-16 ***
SuburbBrunswick	870.8069	56.3822	15.45	<2e-16 ***
SuburbHawthorn	1160.2793	61.6807	18.81	<2e-16 ***

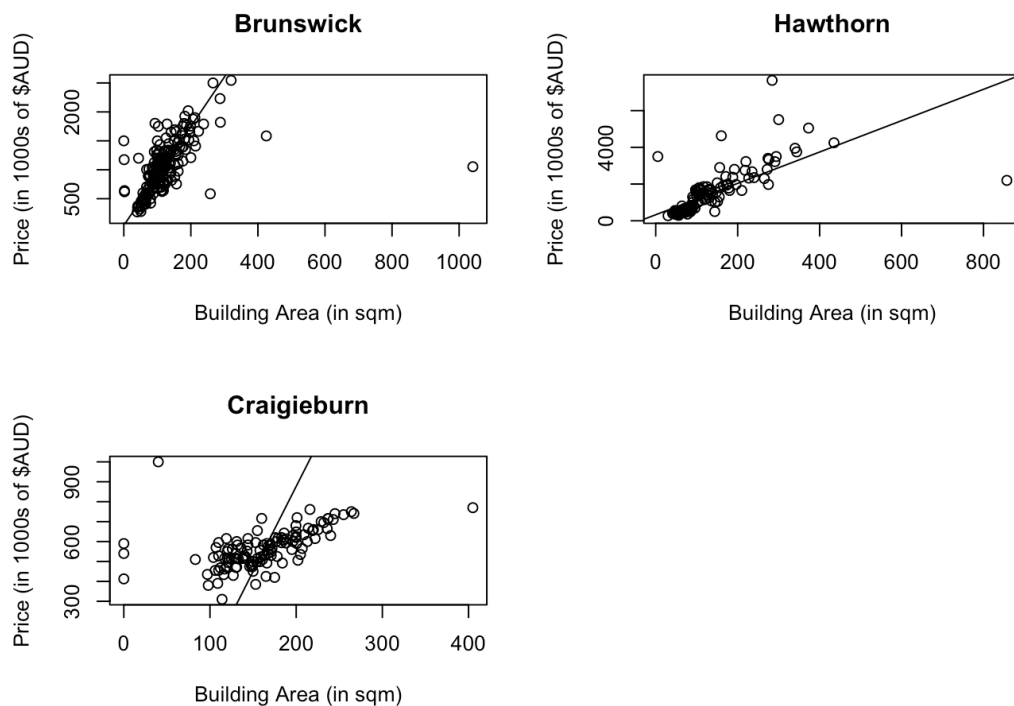
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 442.6 on 396 degrees of freedom

(81 observations deleted due to missingness)

Multiple R-squared: 0.5984, Adjusted R-squared: 0.5954

F-statistic: 196.7 on 3 and 396 DF, p-value: < 2.2e-16



Removing the outliers does improve the fit of the model as the overall residual standard error has decreased and the goodness of fit metrics have increased to around 0.6. Also visually we can see an improved fit for Hawthorn and Brunswick. Although there still is a poor fit for Craigieburn

2.6 Prediction

Predict the price of a house in Hawthorn with 2 car spaces and 100 sqm in building area. What is the 95% confidence interval of your prediction value?

Solution

▼ Code

```
predict(lmfit4, data.frame(Suburb = "Hawthorn", BuildingAr
```

```
      fit      lwr      upr
1 1514653 1393868 1635439
```

If using the `lmfit4` model, then the predicted value for an average property with those features would be \$1,514,653.