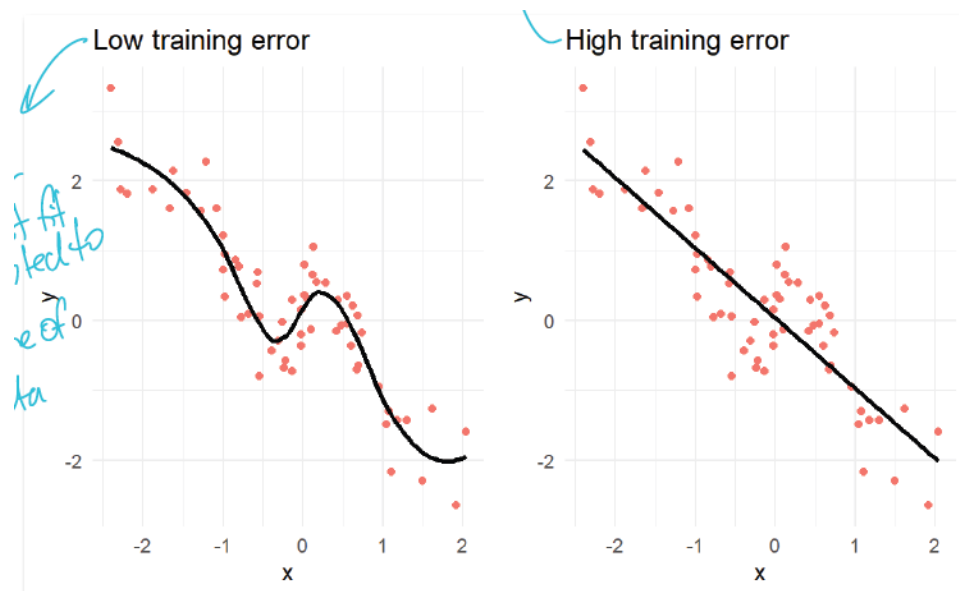



# Week03 - Summary

## Training vs Testing

- **Training error** is the performance metric applied to the observations used to train the model
- **Test error** is the average error when applying a model to predict the response on new (test) observations that were not used in the training of the model
  - Training error can underestimate the test error





### Bias-Variance Tradeoff

- We observe that the more complex we make our model, the better it seems to perform, i.e. the prediction error gets lower in the training sample
- However, in the test sample, we must be cautious not to overfit the model. Hence we must find the right balance between bias and variance

### **Common test set approach**

- Randomly divide the available set of samples into two:
  1. Training set
  2. Test set
- The model:
  - Is fitted on the training set
  - Can be used to predict the responses in the test set

- Quality of the predictive performance uses the test set with:
  - Mean square error (MSE) for a quantitative response
  - Misclassification rate for a qualitative (categorical) response

### Drawbacks of the test set approach

- The estimate of the test error can be highly variable
  - Depends on assignment of observations to the training and test sets
- Only a subset of the observations are used to:
  - Fit the model
  - Assess the model
- Test set error may overestimate the test error compared to a model fit on the entire data set

## Cross-Validation

### K-Fold Cross Validation

- Widely used approach for estimating test error
  - Estimates can be used to select best model and to give an idea of the test error of the final chosen model
- Idea: randomly divide the data into  $K$  equal-sized parts
  - Leave out part  $k$ , fit the model to the other  $K - 1$  parts (combined), and then obtain predictions for the left-out  $k$ th part
- This is done in turn for each point  $k = 1, 2, \dots, K$  and then the results are combined

### Example: 5-Fold

1	2	3	4	5	6	7	8	9	10	11	12	13	14	16	17	18	19	20
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----	----	----	----	----

7	6	3	13	1	4	2	20	9	5	15	10	8	19	11	18	16	14	12	17
7	6	3	13	1	4	2	20	9	5	15	10	8	19	11	18	16	14	12	17
7	6	3	13	1	4	2	20	9	5	15	10	8	19	11	18	16	14	12	17
7	6	3	13	1	4	2	20	9	5	15	10	8	19	11	18	16	14	12	17
7	6	3	13	1	4	2	20	9	5	15	10	8	19	11	18	16	14	12	17

These are folds  
h = 5.

## Cross-Validation Process in Regression

- Define  $K$  folds as  $C_1, C_2, \dots, C_K$ , where:
  - $C_k$  denotes the indices of the observations in fold  $k$  ✓
  - There are  $n_k$  observations in fold  $k$  ✓
- Compute error metric for each fold
  - Mean square error  $MSE_k = \sum_{i \in C_k} (y_i - \hat{y}_i)^2 / n_k$
  - $\hat{y}_i$  are the predictions on the test fold,  $k$
- Aggregate over all folds
  - Overall error =  $\sum_k MSE_k / K$

## Cross-Validation for Classification Problems

- Same as regression but change the assessment metric
- Compute the accuracy for each fold by calculating:
  - $A_k$ : the accuracy of the classifier in fold  $k$ 
    - $A_k = \frac{1}{n_k} \sum_{i \in C_k} \mathbb{1}_{\{\hat{y}_i = y_i\}}$
    - $n$ : the total number of observations in the data set
    - $n_k$ : the number of observations in the belonging to class  $k$
- Aggregate over all folds
  - Overall accuracy =  $\frac{1}{K} \sum_k A_k$

## Repeated Cross-Validation Properties

Repeated CV provides a less-biased CV error estimate

- Repeated CV also gives you the variance of the CV error
- However, it comes with a computational cost
- Implemented in the `caret` package in R

## Example of CV Procedure

- Consider a high-dimensional data set
  - All entirely numeric
  - Need dimension reduction to proceed
- You decide to use the following CV procedure:
  1. Compute correlation matrix, select the top 50 variables that have the highest correlation with the response
  2. Use these 50 variables as features and perform K-fold cross-validation
- Variable selection performed once using both the training and test data sets
- Information can leak from the test to the training set
- Hence, the CV error estimate is likely to be biased
- Ideally, refrain from using the test data in any way in the training step

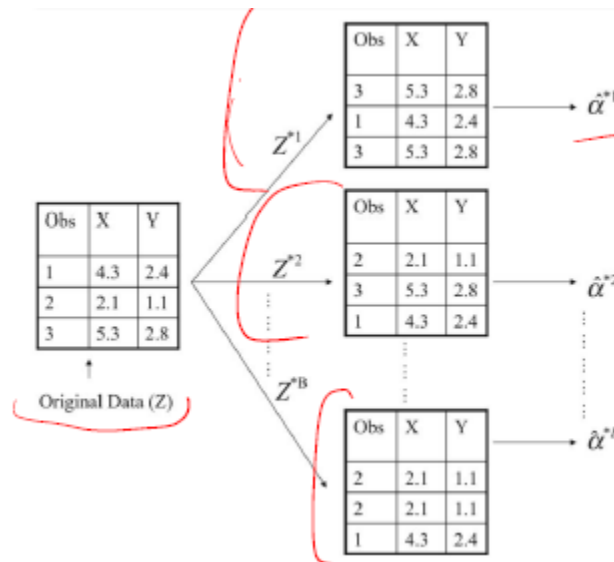
## Corrected CV Procedure

- Split the data set into K folds
- For each  $k = 1, 2, \dots, K$ 
  - Determine the variables that correlate the best with the response using all the data except the data in fold  $k$

- Train your model using the selected variables above
- Run your classification algorithm and record accuracy against the test set

## Bootstrap

- Essentially sampling with replacement



## Simple Investment Example

- Goal: invest a fixed sum of money in two financial assets
  - That yield returns of X and Y where X and Y are random quantities
- Problem: need to decide allocation in each asset
  - Invest fraction  $\alpha$  of our wealth in X and  $(1 - \alpha)$  in Y
- Criterion: minimise the total risk of the investment
  - Mathematically involves minimising  $Var(\alpha X + (1 - \alpha)Y)$
  - The solution (via calculus) is:

$$\alpha = \frac{\sigma_Y^2 - \sigma_{XY}}{\sigma_X^2 + \sigma_Y^2 - 2\sigma_{XY}}$$

Where  $\sigma_X^2 = \text{Var}(X)$ ,  $\sigma_Y^2 = \text{Var}(Y)$  and  $\sigma_{XY} = \text{Cov}(X, Y)$

Proof:

$$\text{Var}(\alpha X + (1-\alpha)Y) = \alpha^2 \text{Var}(X) + (1-\alpha)^2 \text{Var}(Y) + 2\alpha(1-\alpha) \text{Cov}(X, Y)$$

To minimise this expression, we take the derivative with respect to  $\alpha$  and set it to 0:

$$\frac{\partial}{\partial \alpha} [\alpha^2 \text{Var}(X) + (1-\alpha)^2 \text{Var}(Y) + 2\alpha(1-\alpha) \text{Cov}(X, Y)] = 0$$

Simplifying, we get:

$$2\alpha \text{Var}(X) - 2(1-\alpha) \text{Var}(Y) + 2 \text{Cov}(X, Y) = 0$$

Solving for  $\alpha$ , we get:

$$\alpha = \frac{\text{Var}(Y) - \text{Cov}(X, Y)}{\text{Var}(X) + \text{Var}(Y) - 2 \text{Cov}(X, Y)}$$



$$\alpha = \frac{\sigma_Y^2 - \sigma_{XY}}{\sigma_X^2 + \sigma_Y^2 - 2\sigma_{XY}}, \text{ where } \sigma_X^2 = \text{Var}(X), \sigma_Y^2 = \text{Var}(Y), \sigma_{XY} = \text{Cov}(X, Y)$$

- Suppose that X and Y can be sampled from the population repeatedly
- Can estimate the standard deviation of  $\alpha_{\text{hat}}$ 
  - Paired observations (X, Y) can be repeatedly simulated, say 100 pairs to get a single estimate of  $\alpha$ ; repeat this process to get 1000 estimates for  $\alpha$

## Regression Trees

- Segments the players into three regions.

1.  $R_1 = \{X | \text{Years} < 4.5\}$  ✓
2.  $R_2 = \{X | \text{Years} \geq 4.5, \text{Hits} < 117.5\}$  ✓
3.  $R_3 = \{X | \text{Years} \geq 4.5, \text{Hits} \geq 117.5\}$  ✓



### Details on tree building process

The goal is to find boxes  $R_1, \dots, R_J$  that minimises the residual sum of squares (RSS), given by:

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

Where  $\hat{y}_{R_j}$  is the mean response for the training observations within the  $j^{th}$  box

### Tree building with recursive binary splitting

- Take a top-down, **greedy** approach that is known as **recursive binary splitting**
  - Top-down:** begins at the top of the tree
    - Then successively splits the predictor space
    - Each split is indicated via two new branches further down on the tree
  - Greedy:** at each step: the best split is made at that particular step
    - Rather than looking ahead and picking a split will lead to a better tree in some feature step




## Using regression trees for prediction


- Partition the predictor space
  - The set of possible values for  $X_1, X_2, \dots, X_p$  into  $J$  distinct and non-overlapping regions,  $R_1, R_2, \dots, R_J$
- Mean value with each region used for prediction
  - If observation in region  $R_J$ , mean of training observations  $R_J$  used for prediction



## Decision Trees

- What if baseball player salary was categorical (ordinal)? **Salary** is colour-coded





## Decision Trees for Classification

- Very similar to a regression tree
  - Exception being the production of a qualitative response rather than a quantitative one
- For a classification tree
  - Inspect the region that the observation belongs and predict the most commonly occurring class in that region



## Gini Index: Tree building with qualitative target

- Still recursive binary splitting to grow a classification tree
  - Can't use RSS as criterion for binary splits
- Alternative measure such as Gini Index is used instead
- The Gini Index is defined by:

$$G = \sum_j \sum_{k=1}^K \hat{p}_{jk} (1 - \hat{p}_{jk})$$


$\hat{p}_{jk}$  represents the proportion of training observations in the  $j^{th}$  region that are from the  $k^{th}$  class

### Gini Index

- Is a measure of total variance across the K classes
- Gini index is referred to as a measure of node **purity**
  - A small value indicates that a node contains predominantly observations from a single class

### Cross-Entropy: An alternative to Gini

- $D = - \sum_m \sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}$
- It turns out that the Gini index and the cross-entropy are very similar numerically



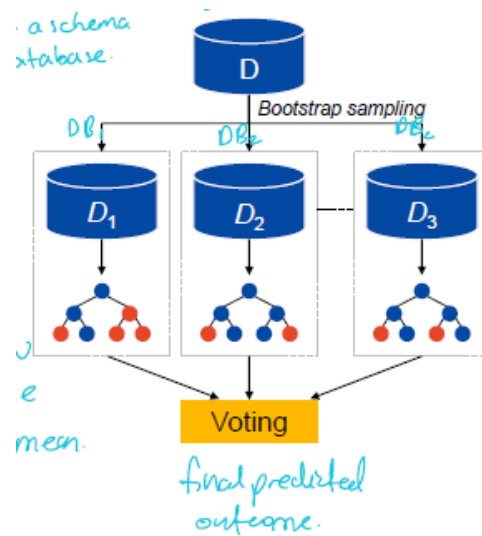
## Ensemble Methods with Bootstrapping

### Bagging (Bootstrap Aggregation)

- Bootstrap aggregation, or bagging
  - General purpose procedure for reducing the variance of a statistical learning method
  - Useful and frequently used in the context of decision trees
- Recall that given a set of  $n$  observations  $Z_1, \dots, Z_n$  each with a variance of  $\sigma^2$ 
  - The variance of the mean  $\bar{Z}$  given by  $\sigma^2/n$
  - In other words, averaging a set of observations reduces variance
- Usually not possible to have access to multiple training sets
  - Using **bootstrapping** instead to create multiple training sets
- Use bootstrapping; can generate  $B$  different bootstrapped training data sets
  - Train the method of the  $b^{th}$  bootstrapping set in order to obtain  $\hat{f}_b^*(x)$ , the prediction at a point
- Average all the observations to obtain:

$$\hat{f}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b^*(x)$$

- This is called bagging



## Out-of-Bag Error Estimation

- Test error of a bagged model is easy
- Each bagged trees using bootstrapped data
  - Only around two-thirds of the original observations are used
  - The remaining one-third of the observations not used
  - Referred to as the out-of-bag (OOB) observations
- Enables prediction of the response for the  $i$ th observation using each of the trees in which that observation was OOB
- This will yield around  $B/3$  predictions for the  $i$ th observation, which we average
- The estimate is essentially the leave one out (LOO) cross-validation error for bagging, if  $B$  is large

## From Bagging to Random Forest

- Random forests (sometimes) provide an improvement over bagged trees
  - A small change decorrelates the trees
  - Reduces the variance when we average the trees
- Random forest approach
  - Create decision trees on bootstrapped training samples
  - When tree fitting, each time a split in a tree is considered
    - A random selection of  $m$  predictors is chosen as split candidates from the full set of  $p$  predictors
  - A fresh selection of  $m$  predictors is taken at each split

The fundamental difference between bagging and random forest is that in Random forests, only a subset of features are selected at random out of the total and the best split feature from the subset is used to split each node in a tree, unlike in bagging where all features are considered for splitting a node.

## Ensemble Methods with Boosting

- Instead of computing a single tree once
- Learn from the mistakes of the current tree
- Take a sample of decision trees into account
- Make final prediction model using the aggregated result from an ensemble of trees

### Boosting

- Bootstrap aggregation combines the results of separately grown trees
- Boosting similar except that the trees are grown *sequentially*

### Boosting for Regression Trees

1. Set  $\hat{f}(x) = 0$  and  $r_i = y_i$  for all  $i$  in the training set
  - Here  $r_i$  denotes the  $i^{\text{th}}$  residual and  $y_i$  the outcome
2. For  $b = 1, 2, \dots, B$ 
  - Fit a tree  $\hat{f}_b$  with  $d$  splits ( $d + 1$  terminal nodes) to the new training data  $(X, r)$ 
    - That is  $r$  is the new response value ✓
  - Update  $\hat{f}$  by adding in a shrunk version of the new tree
    - $\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}_b(x)$  ✓
  - Update the residuals
    - $r_i \leftarrow r_i - \lambda \hat{f}_b(x_i)$

3. Compute the final model

$$\hat{f}(x) = \sum_{i=1}^B \lambda \hat{f}_b(x)$$

scaling factor to incorporate many trees

## Parameters to tune in Boosting

- Number of trees  $B$
- The shrinkage parameter  $\lambda$  a small positive number
  - Typical values are between 0.01 and 0.001
- Number of split  $d$  in each tree
  - Controls the complexity of the boosted ensemble
  - If  $d = 1$ , then the tree is just a stump; this actually usually works quite well

## Bagging vs Boosting

- Both ensemble methods get  $N$  learners from 1 learner
  - Built independently for bagging
  - Built sequentially for boosting

- Trees built-in
  - Boosting are weak learners (sometimes just a stump)
  - Random forest have higher complexity
- Number of parameters to tune
  - Low in random forest
  - High in boosting (depending on which variation you are using)

### Another Boosting Algorithm: Adaboost

- Basic idea
  - At each iteration, reweigh the data to place more weight in data points that the classifier got wrong
- Combine all the weak classifiers by taking a weighted combination
- Put more weight on the weak classifiers with the higher accuracies

